

Next-Generation AI Compiler Survey

Next-Generation AI Compiler Survey

Predicting the agentic compiler (~2027–28 and ~5 years): software-stack reshape and HW–SW codesign

Living survey export

Generated: 2026-08-31

Primary narrative: docs/SURVEY.md (single reading path §0–§9) · Evidence: reference/

North star. Agents own semantic search, orchestration, and artifact synthesis. Compilers own lowering, legality, measurement, and fallback. Hardware codesign enters only through kernels, IR, tests, and profilers — not autonomous tape-out.

Survey narrative

Last updated: 2026-08-31 (skill-scope fold + Beamer T10/C7/★; goal-align)

Evidence store: [../reference/README.md](#) → publications · products · repos

Status: [../STATUS.md](#)

How to read (smooth path).

1. §0 north star + vocabulary → §1–§1b trends → §2–§3 mechanisms → §4 gaps → §5 prediction (architecture, roadmap, stack, commercial, **technical techniques** §5.8).
2. When sources disagree → §6 **Conflicts**. Claim IDs → §7. System snapshot → §8.
3. Digests / SKUs / forges stay under [reference/](#) so the narrative does not become a catalog.
4. Maintainers: §9 add-source loop; python3 scripts/validate_survey.py.

One-page success check: (1) Predicted agentic compiler? → §5.1 / §5.5. (2) Four jobs + stack? → §5.1 / §5.6. (3) Layers + e2e-seeking + **when merge/replace?** → §5.1.1–§5.1.4. (4) Evidence vs noise? → Tier A vs C + §6. (5) Commercial blockers? → §5.7. (6) Which techniques to enhance for the roadmap? → §5.8.

0. North star and vocabulary

0.1 Primary goal


Primary goal: Predict the **next-generation agentic compiler** (architecture + process through ~2027–28 and ~5 years), including how it reshapes the **software stack** and **HW–SW codesign** — without drifting into general EDA.

Everything else (papers, GitHub/Gerrit, commercial SKUs, forums, ASIC bring-up studies) is **evidence** for that prediction—not a catalog for its own sake. When sources disagree, they go in §6 [Conflicts](#) rather than being silently averaged.

Executive verdict. Compilation is shifting from **fixed pass pipelines** + **black-box autotuning** toward **hybrid LLM–compiler loops**. Empirically, the winning pattern is:

Agents own semantic search, orchestration, and artifact synthesis. Compilers own lowering, legality, measurement, and fallback.

The compiler (top of the picture). This is the next-generation agentic compiler. Papers, SKUs, and the rest of this document are **evidence** for that compiler — not a second goal. L1–L7 (§5.1.2) are the bands this loop talks *to*, not extra top-level boxes.



Piece	Role
Goal	Agentic compiler: agents own search; compilers still own legality and codegen.
Control plane	Orchestrates search, admit, freeze, rollback.
Typed kernel	The program the agent is allowed to edit (tiles, roles, barriers, layouts).
Lowering	Deterministic codegen. Not an LLM.
Serve	Load the frozen binary. No agent on the hot path.

Localized reject (where it failed) is a **control-plane edge** back to “try the next kernel,” not a license to rewrite the checker mid-walk. The typed kernel **compiler** may still evolve *across* specialize jobs (new verifier rule / primitive, then re-pin) — that is harness/dialect feedback, not LLM-as-opt. Detail: §5.1.3–§5.1.4, T5 in §5.8.

Agents reshape the **control plane** more than they replace the **data plane**. That control plane is predicted to become **e2e-optimal-seeking** under a product fitness (F) (joint search across multi-band lowers — §5.1.2–§5.1.3); **soft merge of the optimizer (M1) ≠ hard replace of the compiler (M3)** (§5.1.4). A fourth job — **accelerator bring-up / codesign feedback** on sim+silicon — is now Tier A evidence (TritonX, KernelEvolve, Zomboss), still centered on kernels/IR/oracles. August 2026 sources (Cake typed schedule IR, Zomboss compile-once mapping contract, GEAK v4 serving A/B, T-LLM Alive2+CBMC) plus **Argus** (compile-time data-flow invariants + SMT) **reinforce** this hybrid lean — they do not move the goal toward LLM-as-opt. **LLM-oriented IR** is the **agent-visible face** of existing bands (summaries, intent/actions, typed mutation surfaces + admit) — **not** a new L-band and not a license to paste LLVM/MLIR into the prompt (§0.2, **LLM4IR**). **TIRx** (TVM Tensor IR next) is that face on the TVM kernel path: FFI + tile primitives at **L4**, still classical lowering. Vendor L4 DSLs digested after the TIRx miss (**TileLang**, **Gluon**, **TLX**, **CuTe DSL**, **FlyDSL**, **ThunderKittens**)

contest C4 (more sinks) and do **not** add an L-llm band; **Event Tensor** is megakernel IR *above* L4 (L6), not a typed agent face. Hybrid lean holds. See §5 (architecture §5.1, roadmap §5.5, stack §5.6, commercial §5.7, techniques §5.8), §6, §4.

Sub-agent substrate (in scope). Multi-agent **workflow compilers**, **AGI compilers** that freeze agent graphs into deployable artifacts, **static analysis of agent DAGs**, and **heterogeneous agent serving** are first-class evidence for how the control plane is built, secured, and productized—not side topics. Digests: **Auto**, **FlowCompile**, **AgentFlow**, **Heterogeneous agentic AI**. Shipping *runtimes* such as **DeepSeek Harness** (dsh: plugin kernel + append-only session log) sit in the same bucket. So do **Agent Skills** (SKILL.md packaging), **skill compilation** (**SIGIL** ★, **SkCC**, **SkVM**, **SkillSmith**), and **SKILL.state** (explicit (t); discard reasoning after a validated update). They package, compile, or persist what the loop proposes. They are **not** compiler oracles, **not** a fifth job, **not** a new L-band, and they do **not** move the hybrid lean. **Do not collapse three “skill” senses — §0.2.**

0.2 Vocabulary and taxonomy

Two meanings of “AI compiler”

Sense	Meaning	Examples
Compilers for AI	Systems that lower neural graphs to accelerators	TVM, XLA/OpenXLA, MLIR dialects, TorchInductor → Triton, IREE, TensorRT
AI for compilers	LLMs/agents that choose passes, rewrite IR, write heuristics/kernels	Meta LLM Compiler, Compiler-R1, Magellan, GEAK, HintPilot

This survey treats **next-gen** as their **merger**: agents on the control plane, compilers on the data plane.

LLM role taxonomy (New Compiler Stack, 2026)

From *The New Compiler Stack: A Survey on the Synergy of LLMs and Compilers* (arXiv:2601.02045):

1. **Selector** — choose among predefined compiler actions or candidates (pass lists, schedule moves, CUDA template families).
2. **Translator** — rewrite source / IR / assembly (highest correctness risk unless gated).
3. **Generator** — synthesize new compiler artifacts (heuristics in C++, kernels, tools, datasets).

Most strong 2025–26 systems are **Selectors or Generators wrapped in hybrid validation**, not free-form Translators alone.

LLM-oriented IR (agent-visible face, not a new band)

Not the same as IRs for compiling LLM inference (L1–L6) or IRs of agent graphs (ADG / workflow compile — T10). Here: what the LLM should **see or emit** so it can search without becoming opt.

Family	LLM sees / emits	What it is for	Examples	Survey slot
Fluency IR	Raw LLVM IR / asm tokens	Pretrain a prior	Meta LLM Compiler; ComPile ; IRCoder (source↔IR, Code-LM)	T7 data; still a <i>prior</i>
Summary / fingerprint	Compact features, not full IR	Fit context; structure over text	AutoPhase; mlirAgent fingerprints; Fast Feedback metrics; IR2Vec	T1 / T3
Intent / action	Strategy, pass list, hints — not rewritten SSA	Decouple <i>what</i> from <i>how</i> opt <i>does it</i>	Pass-list LLMs; HintPilot; CompileIQ ACFs; IntOpt	C3-B
Typed agent-facing IR	Co-designed language agents author	Mutation surface + pre-compile admit	Cake IR; Argus tags + SMT; Zomboss mapping; TIRx (TVM FFI + tile primitives)	T1/T2 ★ (Cake/Argus); TIRx Tier B

Vendor **L4 kernel DSLs** (not a sixth family and not an L-llm): [TileLang](#) (layout inference + LSP), [Gluon](#) / [TLX](#) (Triton lower/extend), [CuTe DSL](#), [FlyDSL](#), [ThunderKittens](#) — agent *sinks* with explicit layouts/FFI/LSP. [Event Tensor](#) is a megakernel IR *above* L4 (L6 fusion), not a typed agent face. | **Translation (glue)** | One classical IR → another | Interop without rewriting passes | [IRIS-14B](#) GIMPLE→LLVM | §4.4 glue; not a band |

Probe. [LLM4IR](#) / [LaMIR](#) (ICML 2025): models parse LLVM syntax but fail CFG edges and instruction-level execution — fluency ≠ compiler skill. **Lean:** do not add an L-llm band; type the control↔data boundary (T1). Cake vs Argus vs **TIRx** vs vendor tile DSLs (TileLang / Gluon / TLX / CuTe DSL / FlyDSL / TK) remain competing **L4 agent surfaces (C4)**, not a universal LLM IR. (TIRx docs also number *agent-search* rungs L1–L4 — those are **not** this survey’s data-plane L1–L7. **TLX** ≠ **TIRx** ≠ **TritorX**.)

Three senses of “skill” (do not collapse)

In this survey “skill” is a **homonym**. Treat them as different objects:

Sense	What it is	Examples	Survey slot
(1) Vendor compile packs	Installable procedures that drive a <i>compiler</i> search + admit	CompileIQ agent-skills ; TRT-LLM Claude skills	Job (a); P1/P9; still C2
(2) Agent Skills spec	Directory + SKILL.md (YAML frontmatter + Markdown) + progressive disclosure	agentskills.io / Anthropic 2025-12	P1 packaging; not T1
** (3) Skill compilation / (t) **	Compile or VM-execute a skill; or keep explicit execution state instead of a growing transcript	SIGIL (SKILL.md → AG-IR → harness); SkCC (SkIR); SkVM ; SkillSmith ; SKILL.state	T10 / P2 substrate

This repo's `.cursor/skills/` is sense (2) for the *survey authors*, not CompileIQ and not SKILL.state. **SKILL.state is in scope as P2**; its CTF *benchmark* is not compiler evidence (§9, hard rule 15). Skill IRs (AG-IR, SkIR) are **not** the shared agent↔compiler contract (T1) and **not** an L-band.

Agent roles in the compile loop

Role	Job	Examples
Advisor / Selector	Rank candidates, label regions, suggest passes	AgentCompile, Meta LLM Compiler, AutoPass
Translator / rewriter	Source/IR/asm rewrite or hint insertion	ACCLAIM, HintPilot, LLM-VeriOpt
Artifact Generator	Heuristics, kernels, MCP tools	Magellan, GEAK, mlirAgent
Orchestrator	Budget, IR level, stop conditions	ACCLAIM guide agent, GEAK directors
Tester / critic	Tests, Alive2, profiles, refine prompts	ACCLAIM test agent, Generative Compilation
Search partner	Propose nodes for MCTS / evolution	Reasoning Compiler, AlphaEvolve
Bring-up / codesign	Coverage → perf on sim+silicon; ISA/IR feedback	TritonX, KernelEvolve, Ascend diagnosis, KForge

Classical AI compiler stack (substrate)

Framework capture	torch.compile / Dynamo, JAX, TF graphs
↓	
Portable HLO	StableHLO (TF/JAX/PyTorch ↔ XLA/IREE)
↓	
MLIR dialects	multi-level lowers, rewrites, vendor dialects
↓	
Schedule / kernels	TVM TIRx / TileLang, Inductor→Triton/Helion/Gluon/TLX, CUDA Tile IR, CuTe DSL, FlyDSL, CUTLASS
↓	
Serving runtime	vLLM, FlashAttention, CUDA Graphs, decode paths
↓	
CPU / legacy IR	LLVM opt pipelines, PGO / AutoFDO, MLGO advisors

Canonical hybrid loop

Same compiler as §0.1 (top of the picture), compressed:

Capture → Analyze regions → Agent proposes a typed kernel
→ Compiler checks (localized reject → next kernel)
→ Classical lowering (cubin / NPU binary)
→ Measure on the serving path → keep or revert
→ Serve the frozen binary (no model in the loop)

Invariant: LLM outputs guide search; they should not silently define unchecked executable behavior. Serve is a freeze, not another agent turn.

1. What's the trend now? (Q1)

1.1 Two stacks converging

Stack	Question it answers	Maturity
Compilers for AI	How do we lower neural graphs to GPUs/TPUs efficiently?	Mature substrate (TVM/XLA/MLIR/Triton); still fragmented
AI for compilers	How do we choose passes, write heuristics/kernels, use feedback?	Rapid 2023–2026 growth; hybrid systems dominate

“Next generation” is not “throw away LLVM/MLIR.” It is **making those stacks agent-addressable**: structured summaries, bounded candidate spaces, tool APIs, and reward oracles.

1.2 Era timeline

Era	Label	What changed
2018–2022	DL compilers mature	TVM/Ansor, XLA, Glow; MLIR born from TF/XLA needs; cost-model autotune
2020–2023	ML-for-compiler RL	Autophase, CompilerGym, MLGO inlining/regalloc; neural policies hard to ship
2023–2024	LLM enters IR	Pass-list LLMs; Meta LLM Compiler foundation models; compiler feedback loops
2025–2026	Agentic hybrid era	Multi-agent tuners, verifier-RL, heuristic synthesis, Triton kernel agents, generative compilation, vendor CompileIQ / GEAK

1.3 Six active trends (detailed)

Trend A — Hybrid guidance, not LLM-as-compiler

Direct LLM IR rewrite is repeatedly fragile. mlirAgent reports frontier models scoring **below identity** on IR transforms. [LLM4IR](#) shows the complementary failure: models parse LLVM syntax but miss CFG edges and instruction-level execution. Systems that succeed constrain the LLM:

- **AgentCompile**: LLM emits advisory metadata only; templates + checks + benchmarks admit CUDA.
- **HintPilot**: LLM inserts compiler-validated pragmas/attributes, not arbitrary rewrites.
- **Meta LLM Compiler / Compiler-R1**: LLM proposes pass sequences; opt applies them.
- **IntOpt / Cake / Argus / TIRx**: intent sequences or typed agent IRs; classical analysis, lowering, SMT, or TIRx pre-benchmark checks admit. Vendor L4 DSLs (TileLang, Gluon, TLX, CuTe DSL, FlyDSL) are the *sinks* those faces sit on — explicit layouts, not free IR rewrite.

Trend B — From RL gyms to LLM agents

CompilerGym exposed LLVM passes as OpenAI Gym environments (Autophase features, IR instruction-count rewards). That lowered the barrier for RL researchers but policies were often opaque and brittle across compiler versions.

2024–26 shifts toward:

1. **Foundation models** pretrained on IR/asm (Meta LLM Compiler: 546B tokens).
2. **Tool-using agents** trained with SFT+RL (Compiler-R1).
3. **Inference-only multi-agent tuners** that avoid heavy offline training (AutoPass).
4. **Program-synthesis of heuristics** that land as ordinary C++ (Magellan), recovering deployability that neural-in-the-compiler lacked.

5. **Control-plane substrate for sub-agents:** compile/freeze agent workflows ([FlowCompile](#), [Auto](#)), analyze agent programs as ADGs ([AgentFlow](#)), place agent stages on hetero serving ([Heterogeneous agentic AI](#)), and — separately — **compile or state-bound skills** ([SIGIL](#), [SkCC](#), [SkVM](#), [SkillSmith](#), [SKILL.state](#); packaging: [Agent Skills spec](#)) so the control plane is not only a growing chat.

Trend C — MLIR + Triton as default substrate

Practical GenAI path today:

```
PyTorch → Dynamo/Inductor → Triton → PTX/CUDA (or CUDA Tile IR)
```

Parallel portable path:

```
Framework → StableHLO → XLA or IREE → device backends
```

MLIR remains the shared engineering substrate even when product branding differs (OpenXLA, Triton internals, vendor AI compilers, CIRCT). Industry critique (Modular/Lattner) notes fragmentation and that open MLIR AI dialects have not always matched CUDA peak for LLM inference—hence Triton, CUTLASS, FlashAttention, and now CUDA Tile / CompileIQ.

Trend D — Kernel agents go industrial

Writing GPU kernels is the bottleneck for new model architectures and non-NVIDIA portability.

- **KernelBench (ICML 2025):** 250 PyTorch tasks; metric `fast_p` = correct **and** faster than baseline. Frontier models still often <20% one-shot; iterative refinement helps.
- **KernelBench-X:** refinement raises correctness but can **hurt** average speedup; many correct kernels still lose to eager; fusion tasks remain hard.
- **GEAK (AMD):** v3 was repo-level multi-DSL (Triton/HIP/FlyDSL). **v4** adds `e2e_workflow`: Amdahl triage on a warm sglang/vLLM server, recursive kernel search only for kernels that move throughput, warm-server A/B + output parity.
- **Cake (NVIDIA · CMU):** agents author a typed **Cake IR** schedule (roles/barriers/tiers); the harness returns localized diagnostics and itself evolves. Clean-start Cake IR beats matched CUDA/PTX search; Kimi Delta Attention is SGLang-validated ([digest](#)).
- **Argus (CausalFlow / HKUST / Stanford et al.):** tile DSL with **tag functions** / **assertions**; compile-time layout-algebra + SMT admit and thread-level counterexamples. Reports 99–104% of hand-optimized assembly TFLOPS on MI300X GEMM/attention/MoE ([digest](#)).
- **TIRx (Apache TVM):** Tensor IR next — hardware-native kernel DSL with scope / layout / tile-primitive dispatch, **TVM FFI** for construct–inspect–mutate, and pre-benchmark well-formed / sync / race / value-sim checks. Expert B200 kernels track libraries (~0.95–1.00×); the *agent loop* is design, not a public p50 ([digest](#)). **Not** a new L-band (L4). **TileLang** is the TVM *tile* DSL above TIRx (layout inference; 2026 **LSP** inlay hints for inferred layouts).
- **Triton faces (Helion / Triton / Gluon / TLX):** [Helion](#) raises Triton; [Gluon](#) lowers it to explicit layouts + warp specialize; [TLX](#) (Meta, production) extends Triton with MIMW warp-group orchestration. KernelEvolve already searches Triton+TLX.

- **Vendor tile DSLs:** **CuTe DSL** (CUTLASS 4 Python; substrate under ★ CuTeGen); **FlyDSL** (AMD Python+MLIR, CuTe-style algebra; GEAK sink); **ThunderKittens** 2.0 (CUDA-embedded tiles; “you or your agent” is packaging, not a public loop).
- **Event Tensor / ETC (MLSys 2026):** events-as-tensors for **dynamic megakernels** *above* L4 (**digest**) — L6 fusion / AOT vs CUDA Graph recapture, **not** an LLM-oriented IR.
- **Meta KernelLLM:** 8B model specialized for PyTorch→Triton; competitive Pass@k vs much larger general models on KernelBench-Triton.
- **AgentCompile:** compiler-bounded CUDA specialization for transformer inference graphs.

Trend E — Verification enters the loop

Correctness strategies, from weakest to strongest (local):

1. Compile + unit tests / LLM-generated tests (ACCLAIM).
2. Numerical checks vs reference path (AgentCompile).
3. Round-trip / recompile checks (LLM Compiler disassembly).
4. Formal semantic equivalence (Alive2 in LLM-VeriOpt rewards; T-LLM Compiler chains Alive2 + **CBMC** on PolyBench/C).

Formal coverage is still mostly **local peephole / IR / C loops**. Whole-program GPU races and FP nondeterminism lack equally strong oracles. Cake adds *pre-compile* schedule/hardware-conformance gates on a kernel IR — still not a serving-race oracle.

Trend F — Compilers broaden their object

- **Generative Compilation:** sealors complete partial Rust so the compiler can diagnose **during** LLM decoding—not only after file completion.
- **Compiler.next:** treat FMware (prompts, agents, free parameters) as a multi-objective search/compile problem for SE 3.0.
- **Anthropic Claude C Compiler:** agent teams *build a compiler* (~100kLoC Rust) capable of Linux kernel builds—adjacent but important: agents as compiler engineers, not only compile-time optimizers.
- **Magellan / AlphaEvolve:** agents rewrite the heuristic C++ inside production compilers (Google reports production inlining usage and XLA experiments).

1.4 Venue map

Community	What to watch
CGO / CC / ASPLOS / PLDI / MLSys	Systems + compilers; Compiler 2.0 Ken Kennedy Award plenary (HPCA/CGO/PPoPP/CC 2026)
NeurIPS / ICML (+ C4ML)	Methods + KernelBench / Reasoning Compiler / Compiler-R1
ACL Findings	HintPilot-style SE/NLP crossover
LLVM Discourse (LLVM ♥ ML workshop)	MLGO, Magellan, agent PR review
Vendor blogs	NVIDIA CUDA Tile/CompileIQ, AMD GEAK v4, Meta KernelLLM/LLM Compiler, DeepMind AlphaEvolve, Modular
DARPA / labs programs	MOCHA (ML + optimization-guided compilers for hetero HW)

1.5 Public vision works (what's out there)

Besides system papers, several **public agendas** shape the “next compiler” debate. Digests live under Surveys & vision in [../reference/publications/INDEX.md](#).

Vision	Axis	Digest
Compiler 2.0 (Amarasinghe; CC'20 → CGO'22 → Ken Kennedy plenary 2026)	Restore high-level → near-peak on accelerators; ML + better abstractions to <i>build/retarget</i> compilers	compiler-2.0-cgo2026 ★ · lineage '22 · '20
MOCHA / Aarno Compiler 2.0 (funded)	LLM rewrite synthesis + eqsat + Rocq; data-frugal cost models; ISA-as-rewrites	compiler-2.0-mocha-aarno ★
New Compiler Stack survey	LLM as Selector / Translator / Generator; hybrid systems win	new-compiler-stack-survey
Compiler.next	Broaden compile object to FMware (prompts, agents, knobs)	compiler-next
MLIR formal theories	Read AI compilation through formal lenses	mlir-formal-theories
Automated kernel generation survey	Kernel-agent landscape in the LLM era	automated-kernel-generation-survey
IEEE Pulse LLM-compilers outlook	Challenges / future direction essay	ieee-pulse-llm-compilers

Not treated as current vision peers here: pre-LLM CACM “Compiler Research: The Next 50 Years” (2008 NSF workshop); Carbon toolchain modernization talks (orthogonal to AI/agent control planes). Industry substrate critiques (e.g. Modular on MLIR fragmentation) stay in Trend C.

1b. Traditional AI compilation vs following trends

This section compares the **classic AI/DL compiler stack** (graph capture → fixed/autotuned lowering → kernels → runtime) with the **2024–2026 agentic / LLM-hybrid trends** summarized above. The point is not “old bad / new good,” but where each side wins and what the hybrid should keep.

What “traditional” means here

Layer	Traditional approach	Representative systems
Front-end	Trace/export graphs; op fusion by rules	TorchDynamo/Inductor, TF/JAX → HLO/StableHLO
Mid-end	Dialect lowers + handwritten passes	MLIR, XLA, TVM Relay/TIR
Schedule / autotune	Cost models + evolutionary / template search	AutoTVM, Ansor, MetaSchedule
Heuristics	Expert C++ thresholds (inline, regalloc, ...)	LLVM -03, MLGO neural advisors (still in-tree)
Kernels	Libraries + expert CUDA/Triton/CUTLASS	cuDNN, FlashAttention, hand Triton
Correctness	Compiler semantics + unit/golden tests	Deterministic passes; limited formal at scale
Serving	Separate runtime stack	vLLM, TensorRT-LLM, CUDA Graphs

Dimension-by-dimension comparison

Dimension	Traditional AI compilation — pros	Traditional — cons	Following trends (LLM/agent hybrid) — pros	Trends — cons / new risks
Correctness model	Decades of pass engineering; miscompile rates low for common paths	Misses intent-level opts; hard to prove GPU/FP whole-program	Gates (Alive2, tests, templates) can keep compiler as source of truth	Untamed LLM rewrite is unsafe; oracles still local
Search / adaptation	Autotune finds strong schedules given enough trials	Sample-inefficient; weak transfer across shapes/HW/versions	Language priors + MCTS/agents cut samples; multi-level creativity	Token cost, nondeterminism, hard-to-replay traces
Heuristic maintenance	Human-readable, reviewable C++	High expert cost; stale vs new HW/models	Magellan-style synthesis of deployable heuristics; MLGO learned advisors	Ownership, regression, supply-chain of agent code
Kernel specialization	Peak performance when experts invest	Does not scale to every new op/HW (esp. non-NVIDIA)	GEAK/KernelLLM/AgentCompile automate explore-measure loops	Correct ≠ fast (KernelBench-X); fusion still hard
Production readiness	Default paths in PyTorch/XLA/LLVM/CUDA	Fragmented stacks; peak often needs vendor libs	Vendor moves (CompileIQ, GEAK, Magellan prod inlining) show traction	Few agent loops are <i>default</i> compile flags yet
Portability	StableHLO / MLIR aim at framework↔compiler portability	Dialects and Triton/Tile/PTX still siloed	Agents can retarget kernels across AMD/NVIDIA in principle	No standard agent IR/tool contract across stacks
Compile object	Graph → binary/kernels	Does not cover prompts, agents, FMware knobs	Compiler.next / generative compilation broaden the object	Early vision; quality gates immature
Human role	Experts write passes/kernels; long review cycles	Bottleneck on rare expertise	Agents draft; humans review/orchestrate	Review capacity & security become the bottleneck
Cost model	Compile-time CPU + autotune GPU hours (known)	Autotune walls for huge spaces	Can reduce search trials	Adds LLM \$ / latency; budgets poorly standardized
Explainability	Pass lists and schedules are inspectable	Why a heuristic fired can still be opaque	Natural-language rationales + traces possible	Rationales may not match true causal opts

Pros of staying traditional (when to *not* force agents)

1. **Stable, regressable defaults** — -O3, Inductor, XLA pipelines are battle-tested across millions of builds.
2. **Deterministic CI** — same IR in → same binary out (modulo known nondeterminism); agent loops break that unless carefully cached.
3. **Clear ownership** — a pass has a maintainer; an agent trajectory often does not.
4. **Peak library kernels** — FlashAttention/CUTLASS still dominate many hot paths without an LLM in the loop.
5. **Formal/tooling maturity** — Alive2, LLVM lit, FileCheck, and vendor validators assume classical artifacts.

Pros of following trends (when agents earn their keep)

1. **Per-program / per-model specialization** beyond what global heuristics allow (pass order, hints, CompileIQ knobs).
2. **Faster exploration** of kernel and schedule spaces with language priors (Reasoning Compiler, GEAK).
3. **Closing the intent gap** — algorithm-level rewrites compilers cannot invent (ACCLAIM), when tests gate them.
4. **Compiler engineering leverage** — synthesize heuristics/passes (Magellan) or review opt PRs with domain tools (LLVM Discourse agent review).
5. **New surfaces** — mid-decode diagnostics (Generative Compilation); FMware multi-objective compile (Compiler.next).

Synthesis: keep the substrate, change the control plane

Traditional strengths to preserve

deterministic lowering · library kernels · CI regressability · formal local checks

Trend strengths to adopt

advisory search · multi-agent orchestration · heuristic synthesis · profile/verifier feedback

Anti-pattern

replace opt/Inductor with unconstrained LLM codegen and hope tests catch it

Recommendation: treat traditional AI compilers as the **data plane** and agent/LLM methods as an optional **control plane** with admit/fallback. Measure success by (a) production default adoption, (b) replayable traces, (c) oracle strength—not by microbench speedups alone.

2. How do agents help AI compilation? (Q2)

2.1 Mechanisms (why it works)

1. **Sample efficiency** — Language priors propose plausible schedules/passes instead of blind evolutionary samples (Reasoning Compiler vs MetaSchedule-style search).
2. **Correctness envelope** — Constraining interventions to hints, templates, verified IR, or pass lists avoids unconstrained codegen failures.

3. **Multi-level creativity** — Compilers miss purpose-level rewrites; LLMs can restructure algorithms when tests gate acceptance (ACCLAIM).
4. **Compiler engineering itself** — Agents write maintainable heuristics/passes (Magellan, mlirAgent) rather than shipping opaque neural nets inside opt.
5. **Feedback densification** — Profilers, compiler diagnostics, Alive2, and Fast Feedback turn sparse rewards into iterative refine loops.

2.2 Closed loop (canonical)

The hybrid loop from §0.2:

```

Capture → Analyze regions → Agent proposes
      → Compiler checks & lowers
      → Verify / test (empirical or formal)
      → Benchmark / select
      → Feedback to orchestrator
      → Fallback if unprofitable

```

Variants:

Variant	Twist
Generative Compilation	Feedback mid-decode via sealors
Reasoning Compiler	LLM proposals expand MCTS nodes
Magellan	Evolutionary mutate of C++ heuristics + Vizier autotune
ACCLAIM	Multi-level budgeted agents + test agent
GEAK	Reflexion-style generate-eval-reflect-optimize on GPU
CompileIQ	Evolutionary search over compiler internal knobs (not source)

2.3 What agents should *not* own

- Unchecked IR/assembly emission as production code without admit gates.
- Silent replacement of numerical kernels without golden or statistical checks.
- Orchestration without reliable tool-calling (ACCLAIM: open models fail on malformed tool calls before code quality).

3. Can agents reshape compilation processes? (Q3)

Short answer: Yes for the **control plane**; no (so far) for wholesale replacement of deterministic lowering.

These reshape claims are **evidence for §5 Future prediction**. The hard limits below—and the gaps in §4—are the **blockers** to that predicted future. Tiered repo/product evidence: [../reference/repos.md](#), [../reference/products.md](#).

3.1 Old → new process map

Legacy process	Agent-reshaped process	Evidence
Fixed -O2/-O3/-Oz pipelines	Per-program pass search with LLM/RL policies	Compiler-R1, AwareCompiler, AutoPass, LLM Compiler
Hand-tuned heuristics in C++	Agents synthesize/evolve heuristic code	Magellan (LLVM inlining/regalloc; XLA)
Black-box evolutionary autotune	Context-aware LLM proposals + structured search	Reasoning Compiler (TVM + MCTS)
Compile → fail → human fix	Compiler feedback during partial generation	Generative Compilation
Experts write Triton/CUDA	Agent generate–eval–reflect–optimize	GEAK, KernelBench agents, KernelLLM
Single IR level optimization	Multi-agent choose abstraction level	ACCLAIM
Compile source → binary	Compile intent / Firmware knobs	Compiler.next (vision)
Generic NVCC heuristics	Per-kernel evolutionary compiler controls	NVIDIA CompileIQ
Humans write the compiler	Agent teams implement a C compiler	Anthropic CCC (engineering experiment)

3.2 Hard limits

1. **Direct IR transformation** by LLMs can underperform identity; search/heuristic synthesis wins more reliably (mlirAgent). **Understanding** of dumped LLVM IR is similarly shallow (LLM4IR: CFG/exec fail).
2. **KernelBench-X**: iterative refine can raise pass rates while average speedup falls; ~46% of correct kernels still slower than eager in their setting.
3. **Tool-calling quality** of open models can break multi-agent compilers before code skill does (ACCLAIM).
4. **Formal verifiers** cover local equivalences; they do not yet bless end-to-end GPU serving stacks.
5. **Cost & reproducibility** of agent compile loops (tokens, nondeterminism, hardware noise) remain open (Compiler.next call-to-action).

3.3 Practical architecture recommendation

Design agent-shaped compilers as:

```
bounded candidate spaces
+ structured IR / region summaries
+ advisory LLM
+ deterministic materialize
+ empirical and/or formal admit
+ fallback
```

Split **offline** compiler-engineering agents (Magellan-style heuristic synthesis) from **online** compile-time agents (HintPilot / AgentCompile / CompileIQ).

4. What’s missing / under-covered? (Q4)

The gaps below are not a separate “wishlist”—they are the **blockers to the §5 predicted future** (agent-addressable data plane, **four** agent jobs (a–d), first-class artifacts, classical defaults until CI proves agents). Coverage is uneven. Each gap spells out **what exists, what is missing, why it blocks that future, and what “done” could look like**. Digests: [../reference/publications/](#). Evidence maps (Tier A/B/C): [../reference/repos.md](#), [../reference/products.md](#).

Gap map (priority snapshot)

#	Gap	Severity now	Who feels it first
4.1	End-to-end production evidence	High	Framework / compiler vendors
4.2	Correctness at scale	High	Anyone shipping kernels or IR rewrites
4.3	Cost & reproducibility of agent loops	High	CI / release engineering
4.4	Cross-stack interoperability	High	Multi-backend platforms
4.5	Hardware-native agent interfaces	Medium–High	Agent + compiler tool builders
4.6	FMware / agent-app compilation	Medium	LLM-app / SE 3.0 platforms
4.7	Training data for compilers	Medium–High	Foundation-model builders
4.8	Human-in-the-loop compiler engineering	Medium	LLVM/XLA maintainer orgs
4.9	Security / supply chain	Medium (rising)	Prod infra & open-source
4.10	Unified benchmarks	High	Research + industry comparison

4.1 End-to-end production evidence

What exists. Many papers report kernel-, pass-, or microbench-level wins: Compiler-R1 on IR instruction count; HintPilot on PolyBench; Reasoning Compiler on selected serving kernels; AgentCompile / GEAK on model or kernel suites; Magellan talks claim production inlining use; NVIDIA CompileIQ quotes $\leq 15\%$ on already-hot GEMM/attention.

What is missing. Few agent or LLM methods are the **default** path in PyTorch Inductor, OpenXLA/XLA, or LLVM trunk for arbitrary user programs. Wins are often:

- opt-in flags or research forks;
- evaluated on curated kernels, not full training/serving stacks;
- hard to attribute (custom GEMV / CUDA Graphs / KV cache vs “LLM guidance”).

Why it blocks progress. Without default-path evidence, orgs cannot justify always-on agent compile cost or review burden. Survey claims risk overstating readiness.

Done looks like. Public A/B on major frameworks (e.g., Inductor default vs agent-specialized decode path) with regression suites, rollout flags, and multi-month stability—similar to how MLGO reported persistent QPS gains after deployment.

4.2 Correctness at scale

What exists.

Oracle	Strength	Weakness
Compiler applies passes (LLM Compiler, Compiler-R1)	Semantics of passes inherited	Cannot invent new transforms safely
Unit / LLM-generated tests (ACCLAIM)	Catches many functional bugs	Under-approximates; flaky tests
Numerical checks vs reference (AgentCompile)	Good for kernels with goldens	Tolerance games; training vs infer numerics
Alive2 / formal (LLM-VeriOpt; T-LLM + CBMC)	Strong local IR / bounded-C equivalence	Peephole/local; not GPU concurrency
Cake pre-compile gates	Localized sync/resource/ISA reject before GPU	NVIDIA schedule IR only; not portable
Argus tags + SMT	Compile-time layout/pairing/separation checks; concrete counterexamples	Path-insensitive; MI300X prototype; agent-drafted invariants
TIRx pre-benchmark checks	Well-formedness, sync validity, race-freedom, value simulation before GPU	Expert kernels today; no public agent-loop settlement
Round-trip disassembly checks	Partial trust signal	Exact-match rates still modest

What is missing. Oracles for:

- **Whole-program** semantic equivalence after multi-level agent rewrites;
- **Floating-point** contracts (reassoc, TF32/BF16, reduction order);
- **GPU races**, async copies, warp divergence, and memory model bugs;
- **Serving-level** behavioral equivalence (sampling, KV layout, speculative decode) — GEAK v4 warm-server A/B + parity is the first *named vendor loop*; still not a shared multi-month default-path oracle.

Why it blocks progress. KernelBench-X already shows correct kernels can be slow; the dual risk is “fast but subtly wrong.” Formal methods that stop at peephole leave the highest-value agent actions (fusion, schedule, algorithm rewrite) under-governed.

Done looks like. Layered admit policy: local formal where possible → differential testing on shape grids → statistical serving checks → staged rollout. Shared open oracles for Triton/CUDA Tile IR, not only LLVM IR.

4.3 Cost & reproducibility of agent compile loops

What exists. Fast Feedback (~10× iterate vs full IR in-loop); CompileIQ produces portable Advanced Controls Files; some papers fix seeds or report sample budgets; Compiler.next explicitly calls for reproducible intent compilation and shared traces.

What is missing.

- Standardized **cost models** (tokens + compile minutes + GPU-hours per % gain);
- **Deterministic replay** of agent decisions across model versions;
- **Cacheable compilation traces** keyed by (IR hash, HW, compiler version, agent policy);
- CI policies for flaky speedups (hardware noise, DVFS, contention).

Why it blocks progress. A compile that costs \$N of LLM calls and cannot be reproduced is not a compiler feature—it is an experiment. Release engineering will reject it.

Done looks like. Trace format + content-addressed cache (pass list / hints / ACF / kernel artifact) checked into build systems; budget SLOs; golden replay tests when upgrading the agent model.

4.4 Cross-stack interoperability

What exists. StableHLO for framework↔compiler graphs; MLIR as a shared *idea*; Triton as a de-facto GPU DSL; vendor bridges (experimental Triton→Tile IR). **LLM-oriented IR families** (vocabulary §0.2) already exist *inside* stacks: fluency dumps (ComPile / Meta LLM Compiler); summaries (AutoPhase, Fast Feedback, mlirAgent fingerprints); intent/actions ([IntOpt](#), pass lists, ACFs, hints); typed agent IRs (Cake, Argus, Zomboss, [TIRx](#)); translation glue ([IRIS-14B](#)). Vendor L4 DSLs now digested as *sinks*: TileLang (+LSP), Gluon, TLX, CuTe DSL, FlyDSL, ThunderKittens — not extra families.

What is missing. A portable contract for **agent-visible** state:

Concept	Needed across stacks	Today
Region / fusion candidate	Common schema	Ad hoc per paper
Constraint set (shapes, aliasing, HW)	Shared vocabulary	Embedded in prompts
Action space (pass, hint, schedule, template)	Enumerated & versioned	Closed per system
Reward / admit result	Structured feedback	Free-text logs
Artifact identity	Hashable outputs	Often lost

Dialects, Triton (Helion/Gluon/TLX), CUDA Tile IR, CuTe DSL, FlyDSL, TileLang/TIRx, PTX, and LLVM IR remain siloed; an agent tuned on one stack rarely transfers.

Why it blocks progress. Every new agent re-implements glue. Multi-backend companies cannot share learnings.

Done looks like. An “agent compile interface” RFC (perhaps on StableHLO or MLIR) defining summaries, actions, and admit records—similar in spirit to how PJRT standardized runtime plugins.

Not the same as one data-plane IR. A shared *agent-visible* contract (**LLM-oriented IR**) can sit above **several** classical sinks (graph / mid-IR / kernel DSL / backend). Collapsing those sinks into one cost-model abstraction — or treating raw LLVM/MLIR text as that contract — is a different — and, for Horizon A, weaker — bet; see §5.1.1 and LLM4IR.

4.5 Hardware-native agent interfaces

What exists. mlirAgent: structural IR fingerprinting, knowledge graphs, MCP tool suites; Compiler-R1 tool calls (instrcount, etc.); GEAK hardware feedback loops; CompileIQ search spaces over NVCC/PTXAS internals ([NVIDIA/CompileIQ](#)); **Cake IR** + localized verifier/cost tools (NVIDIA schedule IR); **Argus** tag/assert + SMT over a layout algebra; **TIRx** IR/utilities over TVM FFI + tile-primitive dispatch; **TileLang LSP** inlay hints (shapes/dtypes/scopes/inferred layouts); **Gluon / TLX** explicit layout and warp-group ops; **Zomboss** typed mapping surface on generated ASIC backends; **SCM-side tool APIs** in Archer / [llvm-harness](#) (verify, difftest, autofix tools) bound to a local llvm-project checkout; Gerrit AI Agent Provider APIs for in-UI chat (generic LLMs unless extended).

What is missing. Standards, not demos:

- Common fingerprint / provenance for pass effects;
- Vendor-neutral MCP (or equivalent) tool schemas for “compile / verify / bench”;
- Safe sandboxes for executing agent-proposed kernels at scale;
- First-class exposure of HW counters to agents without scraping profiler text.

Why it blocks progress. Agents without structured interfaces fall back to brittle prompt-and-pray over opt stdout.

Done looks like. LLVM/MLIR/Triton tool APIs that agents can call with typed I/O; fingerprinting as a supported analysis; conformance tests for tool servers.

4.6 FMware / agent-app compilation

What exists. Compiler.next vision: compile prompts, agent topologies, and free parameters under multi-objective quality gates; generative compilation couples compilers into coding agents; industry agent harnesses (Claude C compiler) stress test construction. Concrete substrate is arriving: [FlowCompile](#) (compile-time optimize structured LLM workflows), [Auto](#) (freeze witnessed-deterministic agent spans into WASM “cognition binaries”), [AgentFlow](#) (Agent Dependency Graphs for static analysis), [heterogeneous agent serving](#) (place dynamic agent graphs across CPU/accelerator tiers), [VibeServe](#) (agentic end-to-end serving-stack synthesis with accuracy/perf judges), and [DeepSeek Harness](#) (open plugin runtime + append-only session log; *runs* a live agent, does not compile/freeze the graph).

Skill compilation is also not empty (2026-08-31). The [Agent Skills](#) spec (`SKILL.md` + progressive disclosure) is the *packaging* those compilers ingest — CompileIQ already follows it ([digest](#)). [SIGIL](#) ★ compiles `SKILL.md` → AG-IR → a fused harness (mandated steps become graph nodes). [SkCC](#) compiles to SkIR and analyzes injection. [SkVM](#) AOT/JIT-compiles skills across models/harnesses (the “LLMs as processors” line is a **metaphor** — do not promote it to “the compiler is an LLM”). [SkillSmith](#) (2605.15215) compiles to a *boundary ABI*, not a unified workflow IR. [SKILL.state](#) is P2 (`(t)`; discard reasoning after a validated update), not a compiler. Index them as **Tier B substrate**. **Do not** treat `SKILL.md` or SkIR as the kernel IR, and **do not** treat a CTF / exploit bench as compile evidence.

What is missing. Mature analogues of DL compilers for FMware:

- Stable IRs for prompt/tool graphs (ADG is a candidate, not yet a shared standard). Skill IRs (AG-IR, SkIR) are **not** that IR and **not** T1;
- Compilation that **fails closed** when quality thresholds miss — SIGIL/SkVM/SkillSmith are **not** a fail-closed **compiler** product CI;
- Interoperability between “prompt/workflow/skill compilers” and classical model compilers;
- Shared traces for community learning (Compiler.next call-to-action #10).
- **Kernel** history as a VCS (`git` + tags), not a chat, not DSH’s log, and not SKILL.state’s `(t)`.

Why it blocks progress. LLM applications (and agentic *compiler* control planes) still tune by hand and folklore while DL graphs enjoy decades of compiler investment. Without workflow-compile + freeze + ADG checks, multi-agent compiler products stay demo-grade.

Done looks like. Reproducible FMware / agent-workflow compile pipelines with gold labels, cost/latency/quality Pareto fronts, static ADG admit, and CI that blocks regressions—parallel to how model zoos ship compiled artifacts today.

4.7 Training data for compilers

What exists. Meta LLM Compiler: 546B tokens of LLVM-IR/assembly + compiler-emulation instruction data; **ComPile** production LLVM IR pile (Rust/Swift/Julia/C++); **KernelBook** (~18k torch↔Triton) fueling KernelLLM / **TritonRL**; **DR Triton** synthetic CSP-DAG scale-up; Compiler-R1 reasoning dataset (~19.6k); CompilerGym workloads; MLGO training corpora (often internal). **IRCoder** / SLTrans is parallel **source↔IR** for Code-LMs (Tier C for this prediction).

What is missing. Large, open, **versioned** corpora for:

- MLIR dialects (linalg, scf, affine, vendor dialects);
- Triton / Tile IR / PTX paired with schedules and performance labels;
- StableHLO graphs with lowering outcomes;
- Negative data (failed compiles, miscompiles, slow kernels) for critics/verifiers.

Why it blocks progress. Without data, Selector/Generator models stay LLVM-centric or overfit KernelBench. Follow-ons to Meta LLM Compiler for modern AI IRs remain sparse.

Done looks like. Public “ImageNet for compilers” 2.0: multi-IR, multi-HW, with licenses cleared for commercial research—building on CompilerGym’s original ambition.

4.8 Human-in-the-loop compiler engineering

What exists. Magellan produces reviewable C++ heuristics; **Archer** ([paper](#), [GitHub](#)) agentially reviews **LLVM GitHub PRs** with Alive2/LLUBI evidence gates; LLVM Discourse threads report similar agent PR review experience; **Gerrit** hosts general AI review plugins ([ai-code-review](#), [ReviewAI](#), [GerritForge provider](#)) used in large-org change workflows; Anthropic CCC emphasizes harnesses; Lattner commentary stresses tests as the real product. See [../reference/repos.md](#) for the SCM map.

What is missing. Process answers, not only models:

- Who **owns** agent-written passes six months later?
- How do code reviews differ when the author is an agent?
- When do humans override agent heuristics after HW refresh?
- How to mix agent draft + expert finalize without review collapse?

Why it blocks progress. Productivity gains reverse into maintenance debt if agent artifacts are unowned. Generic SWE agents underperform on opt review without domain tools—so HITL design is research-worthy.

Done looks like. Documented workflows: agent proposes → automated oracle gate → human review checklist → ownership assignment → regression corpus update. Metrics for review time vs bugs escaped.

4.9 Security / supply chain

What exists. Sparse public discussion: Anthropic author notes unease about unverified deployed code; classical compiler supply-chain concerns (trusting opt, binary provenance); almost no dedicated papers in this survey's catalog on adversarial *agent kernels*. Adjacent: [SkCC](#) compile-time Anti-Skill Injection on untrusted `SKILL.md` (T9 color for **skills**, not CODEOWNERS for admitted kernels). The Agent Skills spec does **not** require oracles or signing.

What is missing.

- Threat models for **malicious or buggy agent kernels** (silent wrong answers, side channels, DoS via pathological schedules);
- Provenance signing for agent-generated heuristics/kernels;
- Sandbox policies for compile-time execution of untrusted proposals;
- SBOM-like records: which model, prompt, tools, and admit gates produced an artifact.

Why it blocks progress. As Magellan/GEAK/CompileIQ-style artifacts enter production binaries, they become part of the trusted computing base.

Done looks like. Security reviews required for agent-admitted artifacts; reproducible builds; allowlists for online agents; red-team suites for kernel agents.

4.10 Unified benchmarks

What exists (fragmented).

Suite	Measures	Blind spot
CompilerGym / Autophase-style	LLVM IR size / pass RL	Not GPU serving
PolyBench / HumanEval-CPP (+ HintPilot)	CPU runtime with hints	Not IR agents or kernels
KernelBench / KernelBench-X / TritonBench	GPU kernel correct+fast	Not full serving graphs
FlashInfer-Bench	Serving-trace kernels → leaderboard → <code>apply()</code> into SGLang/vLLM	FlashInfer-operator families, not IR→fused→full-graph ladder / cost-to-compile
Paper-specific serving kernels	Attention/MoE/MLP slices	Not comparable across papers
Vendor internal suites	Production truth	Closed

What is missing. A shared ladder:

1. CPU IR opt (size + runtime),
2. Single kernels (CUDA/Triton/Tile),
3. Fused regions,
4. Full LLM serving graphs (prefill/decode, batch, long context),
5. Multi-objective (latency, throughput, memory, quality),

with fixed hardware profiles, reference docks, and leaderboards that separate **correctness**, **performance**, and **cost-to-compile**.

Why it blocks progress. Cross-paper “speedup” comparisons are currently nearly meaningless; the field cannot tell if Selector/Generator methods are improving the same problem.

Done looks like. A community benchmark consortium (LLVM ♥ ML + MLSys + vendor tracks) publishing versioned tasks and forbidding cherry-picked single-kernel headlines without the full ladder.

Cross-cutting research agenda (from the gaps)

Technique-shaped prediction (within vs outside the compiler, missing parts, checkpoint map): §5.8.

Theme	Near-term work	Depends on gaps
Admit & fallback standards	Spec for hybrid pipelines	4.1, 4.2, 4.3
Agent compile IR / tools	RFC + MCP schemas; LLM-oriented IR families	4.4, 4.5
Open multi-IR corpora	MLIR/Triton/Tile datasets	4.7, 4.10
Serving-level oracles	Diff tests + statistical checks	4.2, 4.10
Provenance & HITL	Ownership + signing workflows	4.8, 4.9
FMware compile MVP	Quality-gated prompt/agent search	4.6, 4.3

Follow-on questions for an org adopting this

1. Online compile-time agents (HintPilot/AgentCompile/CompileIQ) vs offline heuristic synthesis (Magellan)—which matches your release model?
2. What is the correctness oracle—Alive2, golden kernels, or serving A/B—and who owns false negatives?
3. Which IR is the **agent-visible** contract (LLM-oriented IR: summaries / intent / typed surface) — LLVM IR dump, MLIR dialect, Triton/Helion/Gluon/TLX, StableHLO, CUDA Tile, CuTe DSL, Cake IR, Argus tags, **TIRx**/TileLang?
4. How do you cache and regress agent compile traces across compiler *and* model upgrades?
5. What is the max \$/build or tokens/build you will spend for a median X% win?
6. Who is the named maintainer of each agent-admitted artifact after merge?

5. Future prediction (what next-gen looks like)


Falsifiable sketch for ~2027–2028, conditioned on conflicts in §6.

5.1 Architecture


The **top of the picture** is the agentic compiler in §0.1: control plane proposes and admits; typed kernel + classical lowering produce a binary; serve loads the freeze. What follows is how that loop sits on **jobs (a)–(d)** and data-plane **L1–L7** — bands are legality/lower surfaces, not a second goal.

Hybrid stack: agents orchestrate; classical compilers execute; silicon feeds the next dialect/ISA RFC. The control plane is itself becoming a compile target (workflow IR → analyze → freeze → place).

HUMAN / PRODUCT INTENT
(NL at the edge · policy · budget)





typed tools · admit records · oracles · FSM / plan bounds



GPU / NPU / ASIC
(sim → silicon)

VCS artifacts
ACF · kernels ·
heuristics ·
memory · traces

Serving runtime
specialize hot paths;
freeze for replay



1. **Compiler becomes agent-addressable**, not agent-replaced — **LLM-oriented IR** (summaries, intent/actions, typed mutation surfaces) + tool APIs + admit/fallback (mlirAgent: free IR rewrite loses to identity; [LLM4IR](#): dumped LLVM IR fails CFG/exec).
2. **Four agent jobs stick**: (a) online specialization, (b) offline heuristic/pass synthesis, (c) compiler engineering & review, (d) accelerator bring-up / codesign feedback.
3. **New first-class artifacts**: ACFs, evolved C++ heuristics, verified kernels, optimization memory, bring-up corpora, replayable traces — **and** frozen agent workflows / cognition binaries when the control plane itself is compiled.
4. **Defaults stay classical** until agents win on *distributions* in CI; hot kernels, size-critical apps, and *new ASICs* adopt first.
5. **Will not happen soon**: unconstrained LLM replaces opt/Inductor without oracles (**C6**); autonomous chip tape-out via compiler agents (**C10**).
6. **Data plane stays multi-band** (~L1–L6 today, + maturing fleet **L7**): not one universal cost-model IR. Cluster/power attach as place/objectives/oracles; missing consolidation → **pluggable interfaces**. Detail: [§5.1.1–§5.1.2](#) (A6/S6).
7. **Target e2e optimum as the architecture center**: multi-band lowers stay; **search and admit reshape** around a joint e2e fitness (F) (not per-band greed). Detail: [§5.1.3](#) (A7/S7).
8. **Merge ≠ replace**: e2e search can unify the *control plane* (Horizon A→B) while classical lowering/admit stays; **agents replacing the compiler** (no classical data plane) is **not** the Horizon A–B bet (**C6-B**, [§5.1.4](#)).

Sub-agent / workflow-compile substrate (must consider). The control plane is itself becoming a compile target:

Line	Claim for agentic compilers	Digest
AGI compiler / freeze	Compile agent traces → deployable artifacts; amortize inference	Auto ★ → P23
Workflow compile	Spec/graph → typed workflow configs across accuracy–latency	FlowCompile ★ → P3/P18
Agent static analysis	ADG + typed deps for audit, injection, unsafe tools	AgentFlow → P1/P22
Heterogeneous serving	Place agent stages across NPU/GPU/CPU under SLOs	Heterogeneous agentic AI → P10/P22

Without these, “agentic compiler” collapses to either unconstrained LLM loops or a classical compiler with a chatbot glued on.

5.1.1 How many data-plane abstractions? (one cost model is not enough)

Question. Compilers own the data plane; agents own the control plane. Does the data plane still need **many** IR/DSL levels — or can agents deal with “everything” through **one** abstraction (e.g. an agent that learns a single optimal cost model over all optimization passes)?

Survey lean: keep several data-plane abstractions; do not bet Horizon A on one universal cost-model layer.

Abstraction band (data plane)	What it owns (classical)	Why agents still need it as a <i>sink</i>
Graph / portable HLO	Fusion regions, shapes, framework↔compiler portability	Amdahl ranking, region propose/admit; StableHLO-class contracts
Mid-IR / dialects (MLIR, Inductor FX/graph opts, XLA HLO internals)	Legality-preserving lowers, layout, memory plans	Fingerprints, pass/hint/ACF actions; free rewrite fails here (mlirAgent)
Kernel DSL (Triton / Helion / Gluon / TLX / TileLang / CuTe / FlyDSL / HIP / TIRx / TK ...)	Tile/schedule search surface for peak kernels	Primary agent training + generate-eval loops (C4)
Backend / ISA / device	PTX/SASS/LLVM MC, HW counters, bring-up	Job (d) codesign feedback; oracles are HW-specific (C9/C10)

A fifth band — **CPU/legacy LLVM pipelines** (PGO, MLGO advisors) — remains for size-critical / server CPU paths (job b), even when GPU DSLs dominate AI serving.

Why “one abstraction + universal cost model” fails the prediction (for now).

1. **Different legality and oracles per level.** Alive2-class local honesty on LLVM IR does not certify GPU races or serving equivalence; Triton unit/golden does not replace graph-level fusion admit. A single scalar cost model cannot be the admit gate across these failure modes (§4.2, T2/T6).
2. **Choosing the level is itself the product.** ACCLAIM’s guide agent allocates budget across abstraction levels — evidence that next-gen systems treat **multi-level choice** as a control-plane decision, not a bug to eliminate.
3. **Cost models stay local.** Classical Ansor/MetaSchedule, MLGO advisors, and MOCHA’s data-frugal cost models improve *within* a pass family or rewrite system. They do not historically transfer as one model from graph fusion → Triton tile → regalloc → serving A/B. Agents amplify search *at a level*; they do not dissolve the levels.
4. **“One agent IR” is already contested.** Portable agent *contracts* (summaries, actions, admit records — T1 / §4.4) are desirable; a single executable IR that replaces StableHLO+MLIR+Triton+Tile is not the Horizon A bet (**C4**, S2).
5. **Collapse path that is allowed.** Offline job (b) can *compile away* some online search into shippable heuristics/advisors so users never see the LLM — but the **runtime data plane** still lowers through classical bands with admit/fallback.

Can one model cover *all* passes — including future ones? How large?

Ask	Survey answer
Train one model that <i>rank</i> s candidate moves across many sinks?	Partially yes as a <i>prior</i> / <i>proposer</i> (LLM-class control-plane skill), not as the admit gate. Meta LLM Compiler-scale IR pretrain already shows broad <i>fluency</i> ; mlirAgent shows fluency ≠ safe free rewrite; LLM4IR shows fluency ≠ CFG/execution skill.
Train one model that <i>replaces</i> legality + layered oracles + all pass families?	No for Horizon A — different labels, oracles, and non-stationary HW/ISA (point 1–3 above).
Cover future passes / new ISAs / new cluster collectives?	Not by one frozen train. Future ops need continual / plugin cost models (per-band advisors + new oracle plugins). A universal weights file cannot predict unmeasured HW behaviors without new labels.
How large if someone still tries a “global ranking” prior?	Parameter count is not the bottleneck. Shipping local advisors are often KB–MB (MLGO-class nets / Ansor cost models). A cross-band <i>proposer</i> looks like a code/IR LLM : ~ 7B–70B+ params and 10¹¹–10¹²-token-class corpora (Meta LLM Compiler: hundreds of billions of IR/asm tokens) — still only a prior. Making it an <i>accurate multi-objective cost oracle</i> across L1–L7 would need orders of magnitude more labeled (program, action, HW, energy, serving-SLO) tuples than exist publicly, refreshed every HW/SKU change.
Practical size bet	Keep small local cost models / advisors per band (often ≪1B; proposal priors only) + an e2e search controller (joint/bilevel under (F) — §5.1.3); do not size one mega-cost-model to “eat” L1–L7 or to replace legality.

What agents *can* unify (control plane, not data plane).

Unify	Do not unify
One agent compile interface schema (region, constraints, actions, admit, artifact hash) across sinks	One IR that is both portable HLO and peak kernel DSL
One e2e search controller / orchestrator under product (F) (budget, band choice, stop, freeze) over many tools (M1)	One learned cost that replaces legality + layered oracles; hard replace of the data plane (M3)
One replay/freeze discipline for artifacts	One online LLM loop that silently defines executable behavior

Falsifiers for this lean.

- A production stack ships **default** lowering where a single agent-maintained cost model selects all passes/kernels **without** classical multi-level admit, and holds correctness + p50 gains for months (**would pressure C3/C6 and this subsection**).

- A portable tile/HLO IR becomes the *only* agent training surface *and* peak path, retiring Triton-class DSLs for serious products (**would settle C4 toward one surface** — still not the same as one cost model over all opt passes).

5.1.2 Predicted abstraction inventory — how many layers, for what, and if they do not consolidate

Headline prediction (Horizon A → early B). Expect about **six stable data-plane bands** (plus optional CPU/legacy), not one. Near-future concerns (cluster compilation, power/energy, multi-objective SLOs) mostly attach as **objectives + oracles + placement plugins** on those bands — they do **not** each invent a full parallel IR stack. If industry fails to consolidate a band, the workable path is **pluggable interfaces** (typed tools / dialect plugins / oracle plugins), not “wait for one IR.”

Predicted bands (what they are *for*)

#	Band	Covers today's compiler jobs	Near-future add-ons that land <i>here</i> (not a new stack)
L1	Framework / graph capture	Trace/export, dynamism, region cut, eager↔compile boundary	Speculative / conditional graphs; multi-model / MoE routing graphs
L2	Portable graph IR (StableHLO-class)	Framework↔compiler portability; high-level fusion legality	Cross-framework reuse; first home for sharding annotations that stay portable
L3	Mid-IR / dialects (MLIR, XLA/HLO internals, Inductor graph opts)	Layout, memory planning, pass pipelines, legality-preserving lowers	Auto-sharding <i>implementation</i> ; pipeline/overlap schedules; power/energy-aware pass choice when counters exist; size vs speed tradeoffs
L4	Kernel DSL (Triton / Helion / Gluon / TLX / TileLang / CuTe / FlyDSL / HIP / TIRx / TK ...)	Peak kernels, tile/schedule search (C4)	Power/perf kernels; fused serving kernels; device-family specials
L5	Backend / ISA / device	PTX/SASS/LLVM MC, registers, barriers, bring-up	ISA/dialect RFCs from agent failures (job d); future-device sim sinks
L6	Runtime / serving execution	CUDA Graphs, KV/decode paths, library kernels, freeze-for-serve	Continuous batching interactions; graph-level vs kernel-level admit; replay under serving oracles (T6)
L7*	Fleet / cluster / multi-device (<i>maturing</i>)	Today: often split across L2–L3 (SPMD/sharding) + runtime place	Collective schedules, hetero device placement, multi-node compile+deploy; agent placement policies (see hetero serving digests)
L0*	CPU / legacy LLVM (<i>optional path</i>)	PGO, inlining/regalloc, MLGO advisors, size-critical apps	Remains for non-GPU / host-side / Magellan-class heuristic sinks

*L7 is the main **Horizon A–B growth band**: not always a brand-new IR today, but functionally required once “compilation” means **cluster placement + collectives**, not only single-GPU kernels. L0 stays when the product is not GPU-only.

Rough count to remember.

- **Must-have for AI compilers today:** L1–L6 (**six**).
- **Must-have once cluster/hetero fleets are first-class:** L1–L7 (**seven**).
- **Optional:** L0 for CPU/size paths.

Agents (control plane) sit **above** these bands via tools; they do not replace the bands.

What does *not* need its own IR band

Concern	Better home	Why
Power / energy / carbon	Objective + counters + admit oracle on L3–L6 (sometimes L7 place)	Same program representation; different reward / constraint. Needs HW energy APIs more than a “PowerIR.”
\$/token, compile \$, latency SLOs	Control-plane policy + layered oracles (P4/P10/P18/P23)	Economics are admit/budget decisions, not a new lowering dialect
Safety / compliance / tenancy	Provenance + sandbox + policy on admit records (P11/P13/P21)	Cross-cutting control-plane, not a data-plane rewrite level
LLM-oriented IR (what the agent sees/emits)	Control↔data contract : summaries · intent/actions · typed mutation surface + admit (§0.2, T1)	Same bands; different <i>face</i> . Raw LLVM/MLIR dump is a bad face (LLM4IR). Not a new lowering dialect.
Agent workflow itself	Control-plane substrate (FlowCompile / Auto / AgentFlow)	Compiles the <i>agent graph</i> , not the neural graph (§5.1)

If consolidated layers do not exist — other ways (plugins)

When industry does **not** converge on a clean L2/L4/L7 IR, do **not** wait for unification. Predicted fallback (already the commercial lean for T1 / §4.4–4.5):

Plugin surface	What it standardizes	Analogy / evidence
Agent compile interface (summaries · actions · admit · artifact hash)	How control plane talks to <i>any</i> sink	PJRT-like RFC over MLIR/ StableHLO/Triton tools (§4.4, T1)
Typed tool / MCP-class servers	compile / verify / bench / profile / place with versioned I/O	mlirAgent, Archer, Compiler-R1, CompileIQ skills (§4.5, P1-C)
Dialect / backend plugins	Vendor lowers register as plugins behind mid-IR	MLIR dialect ecosystem; Hexagon-MLIR-class bridges
Oracle plugins	Pluggable legality / golden / Alive2 / serving A/B / energy	Layered admit (§4.2, T2/T6); FlashInfer-Bench-class rungs
Objective / cost plugins	Local cost models or learned advisors <i>per band</i> (not one global)	MLGO advisors, MetaSchedule/ Ansor-style, MOCHA data-frugal costs
Placement / fleet plugins	Multi-device / cluster place+collective policies	Hetero agent serving; SPMD/ sharding passes as optional modules on L2–L3

Survey lean. Prefer **stable bands L1–L6 + maturing L7**, with **plugin interfaces** at the control↔data boundary, over waiting for a single consolidated IR that covers kernels + cluster + power. Consolidation is welcome *inside* a band (e.g. one dominant agent kernel DSL — **C4**); it is not required *across* bands.

Falsifiers for the inventory.

- Production AI stacks routinely ship with ≤ 3 bands and no distinct kernel DSL, yet hold peak + cluster + energy SLOs (**would collapse this inventory**).
- A widely adopted **FleetIR** (or equivalent) becomes mandatory and separate from L2/L3 (**would promote L7 from “maturing/plugin” to full peer band sooner**).
- Energy becomes a first-class IR dialect industry-wide rather than counters+objectives (**would add an L-power band** — not the current lean).


Pointers. Stack reshape: §5.6. Cross-stack + tools: §4.4, §4.5. Commercial contract options: §5.7 P1. Techniques: T1, T2, T5, T6, T10. E2E-optimal-seeking: §5.1.3.

5.1.3 E2E-optimal-seeking architecture

Question. With ~6–7 data-plane bands, per-band greed yields local optima (great L4 kernel, bad fusion; peak tile, worse serving; perfect place, starved kernels). If next-gen AI compilers **care about e2e optimum**, is an architecture that *targets* e2e optimality possible — and if yes, **what does it look like** (how must the stack reshape)?

Complexity bound (not an excuse to stop). A closed-form proof of a unique global optimum over all bands, SKUs, and future ISAs is intractable (joint discrete space, multi-objective (F), noisy HW, incomplete oracles). That bound is **not** “classical multi-pass never guaranteed it, so next-gen need not try.” If e2e optimum is the product goal, the **control plane and admit path must be rebuilt around (F)**; bands stay as legality/lower surfaces, not as independent optimization silos.

Survey lean: yes — an e2e-optimal-seeking architecture is the Horizon target. Not claimed: mathematical unique global optimum. Claimed: reshape so **product fitness (F)** is the only win condition, with **joint (or bilevel) search** across bands and physical e2e measure as admit.



Reshape vs local-friendly hybrid

Local-friendly (avoid)	E2E-optimal-seeking (target)
Bands = independent optimize-then-lower	Bands = legality / lower surfaces; search is joint
Local cost model can “win” a band	Local cost = proposal prior only ; only (F) admits
Orchestrator schedules leftover budget	Orchestrator / joint policy is the optimizer
Microbenchmark as success	Microbench demoted; serving / fleet / energy is success
Train per-pass labels	Train on e2e trajectories (action stacks → (F))
One-pass L1 → L7 greed	Backtracking / bilevel : upper choices scored by rolled-up (F)
Many search surfaces = many optimizers	Prefer fewer search surfaces , many lower surfaces (§5.1.2 plugins OK)

What the architecture looks like (components)

Component	Role
1. Constitution (F)	Sole product score — latency, energy, \$/token, quality, cluster util; often a Pareto / constrained front, never () local costs
2. E2E search controller	Joint or bilevel policy over fuse/layout · tile/schedule · serve · place; grows ACCLAIM-class guide into a true cross-band optimizer (P3/P22 FSM bounds)
3. Learned e2e surrogate + real oracle	Fast surrogate proposes; periodic physical serving/energy measure admits (T6); surrogate never replaces legality
4. Legality plugins stay	Alive2 / golden / race checks remain band-local; they gate proposals, they do not define “win” (§5.1.1)
5. Eqsat / rewrite / MCTS where legal	Structured search inside a band; rank outside by (F) (MOCHA-class + hybrid loops)
6. Freeze under (F)-admit	Ship ACF/kernel/place only when (F) improves vs classical baseline on pinned traces; else fallback (C5/ C6)
7. Offline encode e2e wins	Job (b) distills repeated (F)-winning stacks into heuristics/advisors (Magellan/MLGO; T4)
8. Amdahl-first under (F)	Region cut by e2e impact before deep L4/L5 burn (AutoKernel; GEAK v4 pct_gpu_time × achievable_speedup; P23)

Anti-patterns (fail the e2e-optimum goal)

Anti-pattern	Why
Optimize L1 → L7 greedily once , never revisit	Lower peaks constrain upper choices; place/serve invalidate kernel “wins”
Sum per-band local costs as “global”	Units/oracles differ; sum (F) (§5.1.1)
Report only kernel microbenchmarks	Ignores fusion, batching, collectives, energy (C2)
Unbounded per-band loops without (F) budget	Spend without approaching e2e optimum (P23)
Use “classical never guaranteed optimum” to justify siloed greeds	Complexity bound ≠ product architecture

What “e2e optimum” means here

Target	Meaning
Not claimed	Unique closed-form global optimum over all bands and future HW
Claimed (Horizon A → B)	E2E-optimal-seeking control plane: joint search + physical (F)-admit + freeze/replay; classical default is a competitor inside the same loop , not the score
Ship bar	Non-dominated (or constrained-optimal) on published (F) vs classical baseline under spend SLOs — not “every layer is peak”

Falsifiers for this lean.

- Commercial defaults keep shipping on **per-band microbench** wins without (F)-admit and still satisfy product SLOs for months (**falsifies e2e-as-center**).
- Joint / bilevel e2e controllers **never** beat strong single-band greed + classical lower on shared pinned suites (**would demote the reshape**).
- A single learned policy selects all band actions from (F) **without** legality plugins and holds correctness + p50 for months (**would pressure C3/C6** — stronger than this lean, which keeps legality band-local).

Pointers. Oracles: §4.2, T2/T6. Budgets: P4/P10/P23. Orchestration: P3/P22. Bands/plugins: §5.1.1–§5.1.2. Claims: A7, S7. When merge/replace: §5.1.4.

5.1.4 When do e2e search and layers merge — and when do agents replace the compiler?

Question. If we pursue e2e-optimal-seeking (§5.1.3) across ~L1–L7 (§5.1.2), **when** can everything merge so the AI compiler is **replaced by agents**?

Survey lean: distinguish three merges. Soft/control-plane merge is the Horizon A–B path. Hard replace (agents *are* the compiler, no classical admit/lower) is not predicted through ~2031.

Merge level	What unifies	When (prediction)	Agents “replace” compiler?
M1 — Soft merge (control plane)	One e2e search controller + shared admit/trace bus over many band <i>tools</i>	Horizon A (2027–28) matures; normal by early Horizon B	No — agents own search; compilers still lower/legalize/measure (C6-B , A1)
M2 — Search-surface merge	Fewer <i>agent training / proposal</i> IRs (e.g. dominant kernel DSL or portable contract), while lowers stay multi-band	Contested through 2028 (C4) ; partial consolidation possible ~ 2029–31	No — one search surface ≠ deleting Inductor/XLA/MLIR/runtime
M3 — Hard replace	Default production path is agent-generated lowering without classical admit/fallback underneath	Not claimed through Horizon B (~2029–31) ; only if C6 settlement signal fires	Yes — this is “AI compiler replaced by agents”

TODAY ad-hoc agents on classical bands	HORIZON A M1 soft merge e2e controller over band tools data plane stays	HORIZON B M1 solid + control plane compiled (T10) maybe partial M2 search surfaces	??? (not predicted) M3 hard replace (C6-A settlement) – only if oracles + F-admit make classical lower commercially dead
--	---	---	--

Why M1 can happen without M3

E2e joint search **merges the optimizer**, not the **execution substrate**. Legality, deterministic lower, golden/serving oracles, and fallback are why production stacks trust the path (mlirAgent: free rewrite loses to identity; **C3**). Merging search under (F) *needs* those sinks; deleting them is a different bet.

Preconditions for each level

Level	Must be true first
M1	Typed agent compile interface (T1); layered + serving oracles (T2/T6); replay/freeze (T3); budgeted orchestrator (P22/P23); (F) as sole win (§5.1.3)
M2	Dominant agent training surface <i>and</i> peak path on shared suites (C4 settle); portable summaries still map to multi-band lowers
M3	Production default lowering with no classical admit/fallback holds correctness + p50/p90 (F) for months (C6 settlement); money-grade oracles cover graph → kernel → serve → place failure modes

What users may perceive as “compiler replaced”

Perception	Survey reading
“I only talk to an agent; never invoke opt”	Often M1 + freeze : agent loop in CI, frozen ACF/ kernel at serve — classical data plane still ran under admit
“Vendor shipped an agentic compiler SKU”	Control-plane product over Inductor/XLA/TRT — still hybrid
“No classical compiler binary in the critical path”	M3 — not the Horizon A–B baseline

Tie to conflicts and horizons

- **Horizon A**: M1 begins (jobs a–d, e2e controller early); M3 explicitly does *not* ship (§5.5.1).
- **Horizon B**: M1 solid (compiled control plane, T10); M2 maybe partial; M3 still “will not happen” for end-to-end LLM-as-opt (§5.5.2).
- **C6**: survey bets **B** (hybrid). Flip only on the settlement signal above.

Falsifiers.

- A major AI stack’s **default** serve path has no classical lower/admit and holds (F) for months (**M3 early** — would rewrite A1/A5/C6).
- M1 e2e controllers never beat siloed band agents + classical -03/Inductor on pinned suites (**would demote soft merge**).
- Search-surface consolidation (M2) completes by 2028 across vendors (**would settle C4 early**).

Pointers. C6; A1/A5/A8; roadmap §5.5; techniques T1/T2/T6/T10.

5.2 How agents change the future (process)

Legacy	Agent-changed future	Confidence
Experts hand-write heuristics	Offline agents propose reviewable C++/MLGO features	High (Magellan/MLGO both shipping paths — conflict C1)
Autotune black boxes	Versioned search artifacts (ACF, traces) in VCS	Medium-high (CompileIQ design)
Scarce kernel experts	Multi-agent kernel loops with profiler oracles	Medium (vendor blogs strong; KernelBench-X ceilings — C2)
Human-only opt PR review	Compiler-oracle agents on PRs/Changes	Medium (Archer; generic forge AI is not enough — C7)
Single DSL (CUDA) expertise	Multi-DSL agent skills (Triton/Gluon/TLX/TileLang/CuTe/FlyDSL/HIP/Cake IR/TIRx)	Medium (TRT-LLM agents PR; GEAK v4; Cake; KernelEvolve Triton+TLX — C4)

5.3 What would falsify this prediction

- A major AI stack ships **default** lowering with no classical admit/fallback and sustained correctness.

- Magellan-class synthesis **and** MLGO neural advisors both disappear from production (neither path).
- Kernel agents plateau forever below eager on fusion-heavy suites with no industrial workaround (libraries-only forever).

5.4 Near-term signals conditioning the sketch

From Magellan LLVM Dev Meeting slides ([digest](#)), ACCLAIM ([arXiv:2604.04238](#), [digest](#)), and HW-codesign agents:

Signal	Implication for \$5	Watch
Magellan OSS via OpenEvolve + OSS models	Offline job (b) becomes reproducible outside Google	Public recipes / llvm patches (C1)
Magellan XLA auto-sharding / graph-rewrite green-field	Offline agents invent heuristics where human expertise is thin	End-to-end XLA pipeline eval (slides: WIP)
ACCLAIM multi-level compiler↔LLM cooperation (code)	Online job (a) is orchestration across levels + test admit, not IR replacement	Tool-calling quality; GPU/serving ports
TritonX coverage on MTIA + future-device sim	Job (d) bring-up/codesign enters the agentic compiler	Second-vendor repro (C9)
KernelEvolve hetero NVIDIA/AMD/MTIA perf agents	Production multi-HW control plane	Public traces vs KernelBench-X (C2)
Ascend compiler-grounded Triton diagnosis	Non-CUDA NPUs need IR/pass escalation, not CUDA-pretrained guess	Hierarchy ablations
Helion + CompileIQ ACF path	DSL substrate agents specialize	C4 vs TileLang/CuTe/Gluon/TLX
CuTeGen CuTe generate–test–refine + delayed profiling	Confirms a non-Triton agent training surface (CuTe/CUTLASS lane)	Multi-DSL skills vs Triton-only agents (C4); substrate now digested (CuTe DSL)
CompileIQ agent-skills (AGENTS.md pack + Welch validate)	Vendor packages the online control plane as installable agent skills	Still no public p50/p90 ACF traces (C2)
EmitC-MLGO PoR (June 2026 sync): internal inliner → Android/Fuchsia → Chrome multi-model	Neural-advisor deploy path is advancing, not abandoned	Customer default still unset (C1)
Hexagon-MLIR open Triton→Hexagon NPU stack	Non-CUDA data plane agents can address	Device/SDK-gated oracles (gap 4.4/4.5)
Cake typed schedule IR + evolving harness (arXiv:2608.12629)	Agent-facing L4 IR can beat raw CUDA/PTX search; compiler <i>and</i> kernels co-evolve	Portable Cake-class contract vs Triton/Tile (C4); public tree
Argus data-flow invariants + SMT (arXiv:2604.18616)	Dense compile-time counterexamples let agents reach library-class peak on selected AMD kernels	Public tree; NVIDIA port; KernelBench fast_p / serving A/B (C2)
Zomboss compile-once mapping contract (arXiv:2608.00894)	Job (d) on <i>emerging</i> ASICs works when machine semantics are compiled, not re-prompted	Second-vendor <i>shipping</i> ASIC (C9 still open)
GEAK v4 e2e_workflow (2026-07-08)	Vendor control plane seeks serving (F) (Amdahl + warm A/B + parity), not only kernel microbench	Public p50/p90 + cost-to-compile (C2)
llvm-harness + llvm-bench (arXiv:2603.20075)	Job (c) needs LLVM tools/bench, not generic SWE; true-fix below 22% after expert review	Default oracle-review in llvm-project (C7)

Signal	Implication for §5	Watch
Compiler 2.0 Ken Kennedy plenary + MOCHA (LLM rewrites $\oplus$ eqsat $\oplus$ Rocq; retarget-via-rewrites)	Venue+funding align on verified ML construction / hetero retarget — not free LLM-opt	OSS releases & Year 1–3 evals (digest , MOCHA)

5.5 Roadmap — Horizon A (2027–28) and Horizon B (~2029–31)

Falsifiable sketch conditioned on C1–C10. Architecture target is [§5.1](#); commercialization packaging is [§5.7](#); layer map is [§5.6](#); techniques that accelerate checkpoints are [§5.8](#).

Executive 5-year bet (still agentic-compiler-centric): classical data planes remain; agentic control planes become how orgs survive $O(\text{ops} \times \text{devices} \times \text{generations})$. HW codesign appears as **sim/silicon feedback into IR/ISA/dialects**, not autonomous tape-out (**C10**). New artifacts (ACFs, heuristics, memories, bring-up corpora) sit beside binaries in VCS — and, later, frozen agent workflows when the control plane itself is compiled.

5.5.1 Horizon A — 2027–2028 (near)

What ships

Capability	Predicted state	Leading evidence	Conflicts
Agent-addressable compilers	Tool APIs, structured IR summaries, admit/fallback become normal in LLVM/Inductor/vendor toolchains	ACCLAIM, HintPilot, AgentCompile, mlirAgent (negative free-rewrite)	C3, C6
Online specialization (job a)	Hot kernels/paths use agent or evolutionary search (ACF, hints, Triton/Helion/Cake/Argus refine) in CI for <i>some</i> products — not yet silent default for all builds	CompileIQ, GEAK v4 e2e , AutoKernel, Cake, Argus	C2, C5
Offline heuristic synthesis (job b)	Magellan-class C++ heuristic evolution <i>and</i> MLGO neural advisors both still live (parallel bets)	Magellan; EmitC-MLGO RFC + June 2026 PoR	C1
Engineering agents (job c)	Compiler-oracle PR review (Alive2/opt) in serious LLVM/AI-compiler orgs; generic forge AI stays UX	Archer, llvm-harness / llvm-bench	C7
Bring-up / codesign agents (job d)	Coverage-first ATen/Triton backend generation on sim + silicon becomes standard for <i>new</i> ASICs; compile-once machine contracts (Zomboss-class) for research accelerators	TritonX, KernelEvolve, Ascend hierarchical diagnosis, Zomboss	C9
Verified ML construction (Compiler 2.0 / MOCHA)	Early open releases of LLM → eqsat → formal-admit rewrite / retarget tooling; not yet default production opt	Ken Kennedy plenary 2026; Aarno/MIT/UIUC MOCHA	C3, C6
DSL surface	Triton-family (Triton/Helion/ Gluon/TLX) remains primary agent training surface; TileLang/CuTe DSL/HIP/FlyDSL/ Cake IR/Argus DSL/TIRx /TK force multi-DSL skills	Helion, CompileIQ Helion path, TRT-LLM agents, KForge, CuTeGen, Cake, Argus, TIRx, TLX, FlyDSL	C4

What does *not* ship by 2028

- Unconstrained LLM replaces opt/Inductor end-to-end without classical admit (C6).
- Single “one agent IR” for all vendors (C4 unresolved).
- Kernel agents uniformly beat eager/libraries on fusion-heavy public ladders (C2).
- Agents design *silicon microarchitecture* autonomously (codesign is **feedback to humans/EDA**, not tape-out autopilot) — see Horizon B.

Near-term milestones (watch)


1. Public Magellan/OpenEvolve llvm patches **or** EmitC-MLGO default (C1).
2. CompileIQ/GEAK publish p50/p90 + pinned traces (C2).
3. TritonX-like bring-up reproduced outside Meta (second ASIC vendor) (C9).
4. PyTorch/LLVM release notes list agent/ACF jobs as supported workflows (C5).
5. MOCHA / Compiler 2.0 publishes OSS rewrite+verify evals or retarget demos (program through ~2028).

5.5.2 Horizon B — ~2029–2031 (next ~5 years from 2026)

Still centered on the **agentic compiler** as the product; HW codesign is how that product eats the $O(\text{ops} \times \text{devices} \times \text{gens})$ matrix.

Architecture evolution

How the hybrid stack thickens from ad-hoc agent loops to a **compiled control plane** over a classical data plane. Target stack detail: [§5.1](#).



Read left → right: agents stay on the control plane; the data plane never goes away; what changes is **how compiled, audited, and amortized** the agent graph becomes, and how tightly silicon feedback closes the loop.

Predicted shifts

Theme	5-year outcome	Confidence
Default compile path	Classical lowering remains default; agentic specialize is opt-in then CI-gated default for hot paths	Medium-high
Artifact store	VCS grows first-class ACFs, evolved heuristics, verified kernels, optimization memory, bring-up corpora	High
Heterogeneous serving	Agentic multi-backend kernel/forge loops are how non-NVIDIA fleets stay viable (MTIA/AMD/Intel/NPU)	Medium-high (KernelEvolve, KForge, TritonX)
HW codesign loop	Pre-silicon: agents + sim generate coverage and compiler stress; post-silicon: agents map ISA/IR pain → RFC for next chip / dialect. Humans + EDA still own tape-out	Medium
Verification	Local formal (Alive2-class) + statistical serving oracles + OpInfo-scale suites compose; whole-program GPU/NPU formal still incomplete	Medium
Human role	Experts own oracles, ownership, security, ISA contracts; agents draft kernels/heuristics/backends	High
Will still not happen	Fully autonomous chip + compiler co-generation without human architectural intent; end-to-end LLM-as-opt	High

Codesign-specific roadmap (still agentic-compiler-centric)

Phase	Agentic compiler does	Hardware team gets
Pre-silicon	TritonX-style coverage on QEMU/sim; KernelEvolve-style search under draft ISA docs	Early “can we run NanoGPT/DLRM ops?” signal; IR/dialect bugs
Bring-up	Coverage agents → perf agents ladder	Weeks → hours for ATen/Triton backend skeleton
Steady-state	Online specialize + memory (KernelBlaster) across gens	Portability when L2/TMA/SRAM rules change
Next tape-out	Aggregated failure modes from agent traces (illegal ops, alignment, missing atomics)	Prioritized ISA / memory-system / compiler-pass requests

Non-goal: surveying EDA/RTL LLM tools unless they close the loop into **compiler IR, kernels, or admit oracles**.

5.5.3 Success metrics for this roadmap

1. Can a new ASIC expose an agent-addressable compile/test API and reach >80% ATen coverage via agents within a release cycle? (TritonX bar)
2. Do agentic specialize jobs show **distributional** wins in CI (p50), not only best kernels? (C2)
3. Are Magellan-class and MLGO-class paths both still productive or has one settled? (C1)
4. Does the stack treat traces/ACFs/heuristics as reviewable artifacts with owners? (§4.8–4.9)
5. Can you ship without NL-only contracts — typed tools + replayable admit traces + memory that survives model swap? (§5.7)

Update when CONFLICTS settle or new Tier A codesign evidence lands.

5.6 Stack reshape (SW + HW codesign)

Focus: not “AI software in general,” but how an **agentic compiler** changes layers from framework UX down to silicon feedback. Architecture target: §5.1. Horizons: §5.5. Claim IDs: §7.

5.6.1 Layer map (today → agentic)

Layer	Classical role	Reshape by agentic compiler	Evidence
1. Model / framework	<code>torch.compile</code> , JAX/TF export	Agents consume graphs/regions; Amdahl-rank hot ops; write kernels back into eager/compile path	AutoKernel, Kernel Forge, AgentCompile
2. Kernel DSL	CUDA / Triton / Helion / Gluon / TLX / TileLang / CuTe / FlyDSL / HIP / TIRx / TK	DSL becomes agent training + search surface ; Helion raises, Gluon/TLX lower or extend Triton; TileLang LSP + TIRx FFI; multi-DSL skills required	Helion, GEAK, CompileIQ, KForge, TRT-LLM agents PR, TIRx , TLX , CuTe DSL , FlyDSL
3. Portable IR	StableHLO, MLIR dialects	Must expose fingerprints, tool APIs, legality; free rewrite fails	mlirAgent, StableHLO, MLIR
4. Compiler mid/back	LLVM/XLA/Inductor/NVCC passes	Offline agents evolve heuristics; online agents pick passes/hints/ACFs; MLGO advisors persist	Magellan, MLGO, ACCLAIM, HintPilot, CompileIQ
5. Oracles & profilers	Unit tests, Alive2, NCU	Become admit gates + reward ; federated profilers (MPP) required for hetero HW	Archer, LLM-VeriOpt, KernelEvolve, Ascend diagnosis
6. Artifacts / VCS	Binaries, schedules	ACFs, evolved C++, verified kernels, optimization memory, bring-up corpora	CompileIQ, Magellan, KernelBlaster, TritorX
7. Serving runtime	vLLM, TRT-LLM, custom ads stacks	Agent loops specialize serving kernels; must not break graph-level opts	GEAK, KForge vs TRT-LLM, KernelEvolve
8. Silicon / sim	Manual bring-up, ISA docs	Agents generate backends on sim + silicon ; traces inform next ISA/IR (codesign)	TritorX, KernelEvolve, Ascend NPU paper

5.6.2 Four agent jobs on the stack

- | | | |
|-------------------------|-------------------------|---|
| (a) Online specialize | → layers 1-2-4-5-7 | (CompileIQ, GEAK, AutoKernel, ACCLAIM) |
| (b) Offline evolve | → layer 4 (+ artifacts) | (Magellan / AlphaEvolve) |
| (c) Oracle engineering | → layers 4-5-6 | (Archer, CCC-adjacent) |
| (d) Bring-up / codesign | → layers 2-5-8 | (TritorX, KernelEvolve, Ascend diagnosis, KForge) |

Job **(d)** is the HW-codesign extension: still an **agentic compiler/toolchain** problem (kernels, dialects, tests), not general chip LLM design.

5.6.3 Stack reshape theses (claim IDs)

ID	Thesis	Status
S1	Control plane becomes agentic; data plane stays classical	Supported — see CLAIMS A1
S2	Portability shifts from “write once IR” to “agent + oracle per backend” while IR remains necessary substrate	Contested — C4, C8
S3	New first-class artifacts (ACF/ heuristics/memory/traces) change CI and code review	Supported — A3
S4	Custom ASIC competitiveness increasingly depends on agentic bring-up latency	Supported (industrial) — TritorX/ KernelEvolve; watch second-vendor repro — C9
S5	Profilers and compiler internals move from human IDE tools to agent APIs	Watch — KernelEvolve MPP, Ascend hierarchy
S6	Data plane keeps multiple abstraction bands; agents unify the <i>contract/orchestration</i> , not a single universal cost model over all passes	Supported lean — §5.1.1; watch C3/C4/C6 falsifiers
S7	Control plane is e2e-optimal-seeking : joint search under product (F); bands are lower/ legality, not independent greeds	Supported lean — §5.1.3; watch C2/C6
S8	Soft merge (M1) of e2e search over band tools by Horizon A–B; hard replace (M3) not predicted through ~2031	Supported lean — §5.1.4; C6-B

5.6.4 What *not* to confuse with stack reshape

Lookalike	Why it is weaker for <i>this</i> survey
Generic coding agents on app repos	No compile oracles → Tier C
Pure EDA/RTL LLM without kernel/IR loop	Out of scope unless tied to compiler admit
Vendor SKU lists without agent/oracle APIs	Tier B baselines only

5.7 From prediction to commercial practice — critical problems

§5 predicts a **hybrid** agentic compiler. Shipping that as a product (CompileIQ-class, GEAK-class, Magellan-in-CI, TritorX-style bring-up) forces engineering choices that papers usually skip. Below: **critical problems** → **option sets with pros/cons** → **survey-leaning defaults**. Treat “might be true” options as hypotheses to settle in production, not dogma.

Problem map (architecture → ops → business). P1–P8 = control/data-plane productization; P9–P22 = eval, money, tenancy, legal, versioning, cold start, HITL capacity, flywheel, latency, DR, A/B, compliance, orchestration predictability, attribution—drawn from §4 gaps, CONFLICTS, Tier A digests (KernelEvolve, TritorX, CompileIQ, ACCLAIM, Magellan, GEAK, Archer, CCC), and adjacent agent-production literature (deterministic boundaries, admit-record provenance).

ID	Problem	Lean (one line)
P1	Agent↔compiler contract	Typed tools + structured admit traces; NL only at human edge
P2	Context / memory loss	Scratchpad ≪ dense skills ≪ VCS as product truth
P3	Sub-agents vs monolith	Specialists + orchestrator + shared trace bus; compile/analyze the agent graph (FlowCompile/AgentFlow/Auto)
P4	When agents may run	CI/hot-path → freeze artifacts; not always-on default
P5	Oracles for money	Layered admit + named false-negative owners
P6	Ownership / supply chain	CODEOWNERS + signed provenance + sandbox
P7	Multi-DSL portability	Multi-skill + HW RAG; IR is substrate not sole API
P8	What customers buy	Flag/SKU + frozen artifacts; bring-up = internal/platform
P9	Eval for agents-as-products	Public p50/p90 + private canaries; always show cost-to-compile
P10	Unit economics / pricing	Sell artifacts + CI quotas; not raw token burn
P11	Multi-tenancy / SaaS	Prefer customer-owned CI; cloud needs isolation+residency
P12	Model-provider lock-in	Pluggable models + golden replay; distill when rights exist
P13	Legal / IP of outputs	Customer owns artifacts; opt-in telemetry only
P14	Joint versioning	Pin agent+compiler+HW+model in admit records
P15	Cold start new HW	Coverage SLA then perf SLA (C9)
P16	HITL review bandwidth	Oracle auto-merge narrow; humans CODEOWN rewrites
P17	Trajectory flywheel ownership	Closed for hetero prod; open traces for heuristic research
P18	Interactive latency SLOs	Interactive = lab tier; batch freeze = product
P19	DR / corrupted tree	Allowlisted edits + canary rollback
P20	Production A/B platform	Required before “production default” marketing

ID	Problem	Lean (one line)
P21	Compliance / ISA RAG	Customer-hosted manuals; learn from compiler feedback when needed
P22	Deterministic orchestration	FSM/plan + LLM judgment; not free tool-calling as default
P23	Tokens / inference / model capability	Yes, hard problems — but mostly for <i>online</i> always-on paths; freeze artifacts + route/distill + Amdahl budgets make them manageable

P1 — What is the agent↔compiler contract?

Agents must exchange state with the data plane. The medium of that contract is a product decision.

Option	What it is	Pros	Cons
A. Natural language only	Prompts + free-text tool stdout (opt logs, NCU dumps pasted into chat)	Fast to prototype; matches coding-agent UX	Brittle; non-replayable; model upgrades break CI; no typed admit
B. Structured logs / traces	JSON/protobuf records: IR fingerprint, action enum, oracle result, cost, artifact hash	Replayable; cacheable; audit/SBOM-friendly; CI-regressable	Schema design cost; vendors disagree on fields
C. Typed tool APIs (MCP-class)	compile / verify / bench / profile with typed I/O (mlirAgent, Archer, Compiler-R1 tools)	Clear action space; sandboxes; versionable	Needs toolchain investment; still need a trace store behind tools
D. Hybrid: NL for humans, structured for machines	Engineers chat; agents speak schemas; NL is a <i>view</i> over traces	HITL-friendly without sacrificing replay	Two surfaces to keep consistent

Survey lean. Prefer **C** + **B** as the product contract; use **D** at the human edge. Pure **A** is fine for demos, not for commercial compile paths (\$4.3–4.5). Concrete artifact shapes already exist: CompileIQ **ACFs**, KernelEvolve **MPP** profiler federation, ACCLAIM tool-calling over clang components. The [Agent Skills](#) `SKILL.md` pack (and CompileIQ’s mount of it) is **A/D packaging**, not the admit record. Skill IRs (AG-IR, SkIR) and SkillSmith boundary contracts are **skill-plane** typed tools — still not T1 across MLIR·Triton.

Example (might be true). The durable contract is not “the prompt,” but a **versioned admit record**: {graph_hash, hw_id, compiler_ver, action[], oracle[], artifact_digest, policy_id}. NL rationales are optional commentary attached to that record.

Also commercially. Who governs the schema (vendor vs LLVM/StableHLO RFC)? What is the **fail mode** when the contract breaks (silent wrong admit vs hard fail)? Keep chat UX as a *view* over traces so SaaS does not create two sources of truth.

P2 — Context-window / memory loss across long optimize loops

Kernel and heuristic search run for hours and hundreds of trials. The agent forgets early failures, successful tilings, and HW constraints.

Option	Mechanism	Pros	Cons
A. Stuff the window	Ever-growing chat of logs + code	Simple	Hits context limits; attention dilutes; expensive
B. Dense session memory	Compress trajectory → summary / skill cards / vector store (KernelEvolve skill library; KernelBlaster persistent CUDA KB)	Survives long runs; transferable across tasks	Compression loses edge cases; retrieval can be wrong
C. External artifact memory (VCS)	Check in ACFs, kernels, heuristics, admit traces; agent loads by hash	Durable across jobs/ models; reviewable; SBOM-ready	Needs ownership & review process (§4.8–4.9)
D. Many short-lived sub-agents	Fresh agent per trial/op/ HW; orchestrator holds only pointers	Avoids context rot; parallelizable	Coordination overhead; may rediscover failures without shared store
E. Hybrid: sub-agents + dense + VCS	Orchestrator + specialists; dense working memory; promote winners to VCS	Matches industrial KernelEvolve / ACCLAIM shapes	Highest systems complexity

Survey lean. Commercial practice needs **E**, with a hard rule: **anything that affects a release build must live in C (external artifacts), not only in chat**. Dense memory is for *search acceleration*; VCS memory is for *product truth*.

Adjacent substrate (do not promote to product truth). [DeepSeek Harness](#) is an **append-only** session log — reconstructable, not a kernel VCS. [SKILL.state](#) makes execution state (`_t`) first-class and **discards reasoning** after a validated update — a fourth *runtime* shape besides stuff-the-window / summary-RAG / VCS. [SIGIL](#) / [SkCC](#) / [SkVM](#) / [SkillSmith](#) compile or VM-execute *skills*. None of these is git + tags for admitted kernels.

Example (might be true). Treat the LLM context as a **scratchpad**, dense memory as a **L2 cache of skills**, and git as **durable store**. If the model is swapped, scratchpad dies, L2 may warm-start, git must still rebuild the binary identically.

Also commercially. Multi-tenant skill/RAG isolation; **invalidate** dense memory on compiler/HW upgrades (KernelBlaster across gens); decide whether customer trajectories may train vendor skills (P13/P17).

P3 — Many sub-agents vs one dense-memory agent

Option	Pros	Cons	Best fit
Monolith agent + dense memory	Simpler ops; one policy to eval	Jack-of-all-trades; noisy tools confuse one policy	Small orgs; single DSL
Specialist sub-agents (gen / debug / perf / verify / HW-RAG)	Clear skills; parallel search; mirrors ACCLAIM / GEAK / KernelEvolve	Orchestration bugs; cost multiplies; inconsistent styles	Multi-HW / multi-DSL products
Hierarchy (orchestrator + workers + judge)	Budget control; staged admit	Judge can be wrong; latency	Production CI with spend caps
Compiled / analyzed topology (workflow IR + ADG + freeze)	Topology is a compile object: Pareto configs, static checks, amortize hot spans	Needs IR + adapters; early ecosystem	Any SLA'd multi-agent product

Survey lean. For commercial multi-accelerator stacks, **specialists + orchestrator** win; invest early in a **shared admit/trace bus** so sub-agents do not keep private incompatible memories (ties to P1–P2). Treat the topology itself as compiler-shaped: **FlowCompile**-style offline workflow search, **AgentFlow**-style ADG audit, **Auto**-style freeze of witnessed-deterministic spans.

Also commercially. Budget per specialist (ACCLAIM); tool-calling quality can fail before code skill; support must debug multi-agent traces. Prefer **deterministic orchestration + LLM judgment** (TritorX FSM, GEAK v4 JS workflows) over unbounded free tool-calling for SLA'd products (→ P22).

P4 — Online cost, flaky speedups, and “when may the agent run?”

Option	Pros	Cons
Always-on online agent at compile	Max specialization	\$ and nondeterminism; release engineering rejects
CI / nightly specialize → freeze artifact	Replayable defaults; human review window	Stale vs new shapes; slower iteration
Hot-path trigger only (Amdahl rank, p99 decode)	Spend where it pays (AutoKernel-style)	Trigger policy itself needs ownership
Offline-only (Magellan)	Users see classical -03	Misses per-model serving wins

Survey lean. Commercial default: **CI/nightly or hot-path → freeze ACF/kernel into VCS**; interactive online agents as an opt-in lab mode until budgets and replay are boring (\$4.1, \$4.3, C5).

Also commercially. Publish budget SLOs (\$/build, tokens/%gain, GPU-hours); flaky-speedup CI policy under HW noise; separate **interactive latency tiers** (P18) from batch freeze; who pays for GPU (vendor cloud vs customer cluster) (P10).

P5 — Correctness oracles strong enough for money

Option	Pros	Cons
Unit / golden / OpInfo	Cheap; TritorX-proven for coverage	Misses subtle perf-correct bugs
Numerical tol vs reference	Kernel-friendly	Tolerance games; training≠infer
Local formal (Alive2)	Strong when applicable	Peephole; weak on GPU concurrency
Serving A/B + canaries	Catches product regressions	Slow; expensive; attribution hard
Layered admit (all of the above)	Defense in depth	Process heavy

Survey lean. Layered admit is the only commercial-grade answer (§4.2). Ship with explicit **false-negative owners**.

Also commercially. Cover FP contracts, GPU races, serving equivalence (§4.2)—not only unit tests. Oracle cost can erase single-digit CompileIQ wins (C2). Decide liability when layered admit passes and production still miscompiles. Top layer needs a real A/B platform (P20).

P6 — Ownership, security, and supply chain of agent code

Agent-written heuristics/kernels are still production code.

Option	Pros	Cons
Agent PRs as human-owned	Clear CODEOWNERS	Reviewer bottleneck
Auto-merge under oracle gates	Scale	Oracles incomplete → silent breakage
Signed provenance (model, prompt, tools, admit)	Audit / SBOM-like	New infra
Sandbox + no network for codegen	Reduces exfil / supply risk	Slows tool use

Survey lean. Human CODEOWNER + signed admit provenance + sandbox; auto-merge only for narrow action classes with strong oracles (Archer-style), not free rewrite (C7, §4.8–4.9).

Also commercially. Threat model beyond “sign something”: malicious kernels, prompt/tool injection, DoS schedules (§4.9). CCC-style “not for production” disclaimers signal culture until harness+oracles are boring. Legal/IP assignment of outputs is separate (P13). HITL *capacity* is separate (P16). DR when agents edit trunk (P19).

P7 — Multi-DSL / multi-vendor portability

Option	Pros	Cons
One DSL to rule them (e.g. Triton-only agents)	Focused data	Breaks on TileLang/CuTe/HIP/FlyDSL/Gluon/TLX (C4)
Multi-skill agents + HW RAG	Matches KernelEvolve / KForge / GEAK	Training and eval explode
Lower to portable IR, agents on IR only	Theory-nice	Free IR rewrite weak (mlirAgent); still need device skills

Survey lean. **Multi-skill + HW RAG** commercially; portable IR remains the *substrate*, not the sole agent language (C3, C4).

Also commercially. Cold-start productization for new ASICs (P15); sell **coverage SLA vs perf SLA** separately (C9); eval-matrix cost across DSLs is a P&L item (P9/P10).

P8 — Product packaging: what customers buy

Option	Pros	Cons
Compiler flag / autotune SKU (CompileIQ)	Familiar; ACF portability story	Gains may be single-digit on hot kernels (C2)
Cloud optimize service	Recurring revenue; centralized HW	Data gravity; reproducibility across tenants
Internal platform only (Meta Ranking Engineer Agent)	Fits ads/ASIC TTM	Hard to productize externally
OSS agent harness + paid oracles/HW	Ecosystem	Support burden

Survey lean. Near term: **flag/SKU + frozen artifacts** for external customers; **internal platforms** for ASIC bring-up. Do not sell “chat with your compiler” as the sole SKU.

Also commercially. Pricing/unit economics (P10); multi-tenancy (P11); map SKUs to jobs (a)/(b)/(c)/(d)—CompileIQ flags ≠ TritonX bring-up platforms ≠ AlphaEvolve-style cloud coding agents ([../reference/products.md](#)).

P9 — Eval & benchmarks for agents-as-products

Buyers need **distributions and cost**, not best-kernel blogs (C2, \$4.1, \$4.10).

Option	Pros	Cons
Vendor-private suites only	Matches production truth	Buyers cannot verify
Public ladder only (KernelBench → fusion → serving)	Comparable; CI-friendly	May understate internal wins; maintain cost
Customer-workload certification	Sells on their fleet	Slow sales cycle
Hybrid: public p50/p90 + private canaries	Honest marketing + real SLAs	Two reporting systems

Survey lean. Hybrid; **always report cost-to-compile beside speedup**. Evidence: KernelBench-X ceilings, CompileIQ docs 2–3% on hot kernels, GEAK eval suites, \$5.5 p50/p90 milestone.

P10 — Unit economics & pricing

Token + GPU cost can dominate single-digit gains (\$4.3; ACCLAIM budgets).

Option	Pros	Cons
Per optimize-job / GPU-hour	Aligns with cloud cost	Punishes thorough search
Per frozen artifact (ACF/kernel)	Matches “ship artifacts”	Hard to price before search
Outcome-based (% latency / \$ saved)	Easy ROI story	Attribution wars (P24-class); baseline games
Seat + CI quota	Predictable OpEx	Leaves headroom on huge wins

Survey lean. **Artifact license + CI quota** near term; avoid pure outcome pricing until A/B attribution (P20) exists. Do not sell unbounded token burn.

P11 — Multi-tenancy, SaaS isolation, data gravity

Option	Pros	Cons
Fully managed cloud optimize	Recurring \$; rare HW	Leakage risk; residency; attack surface
Customer VPC / on-prem agent + cloud model API	Data local	Model lock-in
Fully air-gapped local models	Regulated buyers	Weaker models; support cost
Artifact-only: customer CI runs searcher	Least SaaS risk	Harder recurring revenue

Survey lean. **Customer-owned CI** for external compile SKUs; VPC/air-gap for ASIC bring-up with proprietary ISA docs (KernelEvolve-class RAG).

P12 — Model-provider lock-in & swap survivability

Option	Pros	Cons
Single frontier API	Best tool-calling today	Price/quality cliff
Pluggable backends + golden replay CI	Survives swaps	Quality floor = weakest model
Distill specialists on flywheel	Cheap, controllable	Needs data rights (P13/P17)
Offline artifacts only; model in eng CI	Users see deterministic -03	Misses online specialize \$

Survey lean. Pluggable + replay; distill when rights exist (KernelEvolve RL path; Magellan/OpenEvolve OSS signal). Customer default path remains frozen classical artifacts (**C5**).

P13 — Legal / IP ownership of agent-generated artifacts

Option	Pros	Cons
Customer owns outputs; vendor owns weights	Clean procurement	Blocks silent flywheel on customer data
Shared improvement / opt-in telemetry	Improves product	Enterprises often refuse
OSS-license artifacts by default	Ecosystem	Leaks differentiation
Work-for-hire + indemnification SKU	Enterprise-friendly	Expensive legal

Survey lean. Customer owns ACFs/kernels/heuristics; telemetry **opt-in only**; never silent training on customer graphs (\$4.7–4.8; CCC production caution).

P14 — Joint versioning (agent + compiler + HW + model)

Option	Pros	Cons
Content-addressed admit records in VCS	Replayable; SBOM-ready	Schema tax (P1)
Lockstep vendor releases	Simple support matrix	Slow; multi-vendor fleets break
Policy bundles (“optimize profile vN”)	Productizable	Bundle sprawl
Re-optimize on any bump	Always fresh	Cost explosion

Survey lean. Admit records + policy bundles; forbid silent re-optimize without budget SLO (\$4.3 cache keys).

P15 — Cold start for new HW / ISA (bring-up productization)

Option	Pros	Cons
Coverage-first SKU then perf SKU	Matches C9 ladder; sellable milestones	Coverage ≠ serving peak
Perf-only from day one	Clear ROI narrative	Fails ASIC TTM
Sim-first pre-silicon agents	Early codesign feedback	Sim fidelity gaps
Human skeleton + agent fill	Predictable ownership	Slower agent narrative

Survey lean. **Coverage** → **perf** with explicit SLAs; sim-first when pre-silicon (TritonX, KernelEvolve RAG, Ascend diagnosis).

P16 — Human review bandwidth (HITL capacity)

Agents raise draft volume; reviewers become the bottleneck (\$1b, \$4.8, C7).

Option	Pros	Cons
Human owns every merge	Safest	Does not scale
Oracle auto-merge narrow; human for rewrites	Archer-scalable	Oracle holes
Quota reviews/week + prioritization	Capacity planning	Good patches wait
Compiler-oracle review agents as force multiplier	Matches C7-B	Still needs human CODEOWNER

Survey lean. **B + D**; publish review-time vs escaped-bug metrics (\$4.8 done-looks-like).

P17 — Trajectory / flywheel ownership

Option	Pros	Cons
Fully closed trajectories	Moat (Meta/AMD/NVIDIA)	Weak external trust
Public anonymized traces (Compiler.next #10)	Research + smaller vendors	May leak HW limits
Federated / siloed memories	Privacy-preserving	Hard systems
Sell data foundation as SKU	Clear ownership	Few buyers

Survey lean. **Closed** for hetero production agents; **open traces** for CPU/IR heuristic-evolution research (Magellan/OpenEvolve). ieee-pulse: data scarcity is already a field limiter.

P18 — Interactive latency SLOs vs batch optimize

Option	Pros	Cons
Promise interactive “chat optimize”	Attractive UX	Hours-long MCTS breaks support
Lab tier (best-effort) vs product tier (batch freeze)	Honest SLOs	Two products to explain
Hard timeout + partial artifact	Bounded cost	Users get weak results
Offline-only eng tools	Simplest SLA	Misses serving specialize

Survey lean. Lab vs product tiers; product default = batch/CI freeze (P4). TritorX “overnight,” KernelEvolve long search, CompileIQ evolutionary runs are batch-shaped.

P19 — Disaster recovery when agents corrupt trees

Option	Pros	Cons
Branch-only writes + revert bots	Clear blast radius	Slower landing
Allowlisted edit surfaces (EVOLVE-BLOCK)	Magellan pattern	Constrains creativity
Signed release + canary auto-rollback	Production-grade	Needs P20
Immutable artifact store; never live-rewrite trunk heuristics	Strongest safety	Slower iteration

Survey lean. Allowlist + canary rollback for compiler-source agents; branch-only / reversible overlay for repo-level kernel agents (GEAK v4 overlays).

P20 — Production A/B & experimentation platform

Option	Pros	Cons
Serving canaries only	Real regressions	Slow; expensive
Shadow compile + offline replay benches	Cheap gate	Misses serving numerics
Full experimentation platform	Settles C2; funds outcome pricing	Large eng investment
Trust vendor blogs	Zero cost	Not commercial-grade

Survey lean. Shadow/replay as default gate; canaries/full platform **before** any “production default” claim (\$4.1; MLGO persistent-QPS bar).

P21 — Compliance: export, residency, proprietary ISA docs

Option	Pros	Cons
No proprietary manuals in cloud RAG	Lower risk	Weak cold start
Customer-hosted RAG corpora	Residency OK	Integration pain
Certified regions + export classification	Enterprise sales	Slow GTM
Learn from compiler/crash feedback only	TritorX-like	Slower perf ramp

Survey lean. **Customer-hosted manuals** when selling to silicon vendors; else prefer compiler-feedback learning over exporting ISA corpora.

P22 — Deterministic orchestration vs free LLM judgment

Adjacent agent-production work stresses **deterministic boundaries** and moving the LLM out of the hot execution loop; TritorX deliberately uses an FSM rather than free tool-calling. Control-plane substrate papers sharpen the same lean: [AgentFlow](#) ADGs make agent programs analyzable; [FlowCompile](#) pushes config search offline; [heterogeneous agent serving](#) places stages under cost/SLO policies rather than “run everything on the biggest GPU.”

Option	Pros	Cons
Free tool-calling agent	Flexible	Hard to SLA; runaway loops/cost
Plan/FSM + LLM fills bounded slots	Predictable; auditable	Less “agent magic” marketing
Compile-then-execute (LLM offline only)	Aligns hybrid prediction; Auto/FlowCompile path	Needs strong offline jobs (b)/(d) + workflow compile
Soft max-steps + human gate	Practical	Caps autonomy

Survey lean. **FSM/plan + bounded LLM slots** for any SKU with an SLA; free tool-calling stays lab-tier (ties P3/P4/P12; supports **C6-B**). Prefer ADG/static checks before deploy and hetero placement policies for cost.

P23 — Tokens, inference performance, and model capabilities

Question. Will LLM **token burn**, **inference latency/throughput**, and **model capability gaps** block turning the hybrid prediction into commercial practice?

Short answer. Yes — they are first-class commercial problems, especially for always-on online agents. They are **not** show-stoppers for the hybrid bet if products **freeze artifacts**, **Amdahl-budget** search, and **route/distill** models. They *are* show-stoppers for “chat with the compiler on every build” without spend and latency SLOs.

Evidence from this survey's corpus

Dimension	What we see	Digests / sections
Token / \$ cost	Agent loops multiply spend vs one-shot chat; cost poorly standardized (tokens + GPU-hours per % gain); ACCLAIM must <i>distribute budget</i> across levels; AutoKernel warns not to burn tokens on tiny Amdahl slices; ieee-pulse stresses cost/data limits in HPC	§4.3, P4/P10; acclaim, autokernel, ieee-pulse-llm-compilers, Compiler.next CTAs
Inference latency	Kernel/heuristic search is hours-scale (MCTS/evolution, TritorX overnight, CompileIQ evolutionary runs)—not interactive TTFT; multi-agent tool loops add wall-clock even when tokens are cheap	P18; tritorx, kernelevolve, compileiq-*
Model capability — codegen	Frontier one-shot KernelBench often weak (<~20% fast_p); iterative refine helps; specialized small models (KernelLLM 8B) can compete on narrow tasks	kernelbench, kernelllm, Trend D
Model capability — IR rewrite	Free IR transforms can score below identity (mlirAgent); dumped LLVM IR fails CFG/exec (LLM4IR); capability ≠ safe data-plane replacement	mliragent, llm4ir, C3 , A5
Model capability — tools	Open models often fail tool-calling before code quality (ACCLAIM); multi-agent compilers die on malformed tools	acclaim, §3.2
Model capability — sample efficiency	Language priors can cut search vs blind autotune (Reasoning Compiler; AutoPass inference-only); Magellan/AlphaEvolve spend offline then ship classical code	reasoning-compiler, autopass, magellan
Mitigations already shipping	Freeze ACF/kernel into VCS; HW RAG + skills (KernelEvolve); distill/RL specialists on trajectories; Fast Feedback (~10× vs full IR in-loop); offline job (b) so users never pay LLM at -03 time	P2/P4/P12; CompileIQ ACF; Magellan
Control-plane freeze / workflow compile	Auto : compile witnessed-deterministic agent spans → WASM cognition binaries + deopt; FlowCompile : offline Pareto	auto-agi-compiler, flowcompile, agentic-ai-hetero-systems, sigil, skill-state; §4.6

Dimension	What we see	Digests / sections
	configs for sub-agent workflows; hetero placement avoids frontier GPUs for every stage; skill compile (SIGIL / SkVM / SkillSmith) and SKILL.state (_t) cut <i>skill</i> -loop tokens the same way	

Adjacent industry signal (agent production, not compiler-specific). Agentic tasks commonly cost **many**× chatbot tokens (iterative tool use + context re-send—“communication tax”; code-review-like stages dominate token share in agentic SE studies). Prompt caching, model routing (small models for easy steps), and context compaction are becoming mandatory FinOps—not optional polish. This reinforces compiler-agent design: **fewer, higher-value LLM calls** behind oracles beat chatty multi-agent refinement in the hot path—and **compile/freeze the agent graph** when spans are deterministic (Auto), rather than re-paying tokens every run.

Options (how products respond)

Option	Pros	Cons
A. Always-on frontier model at every compile	Max freshness	Token+latency blow up; release eng rejects; capability still needs oracles
B. Batch/CI optimize → freeze artifact (hybrid default)	Amortize tokens over many serves; replayable	Stale vs new shapes; needs re-optimize policy
C. Route + distill (frontier orchestrator, small specialists; KernelEvolve-style RL)	Cuts \$/task and often latency	Needs trajectory rights (P13/P17); tool-calling quality still gates open models
D. Capability firewall (LLM only in bounded actions; classical data plane executes)	Capability gaps don’t miscompile by default	Limits “agent authors the opt” narrative
E. Ignore and hope prices fall	Simple story	Volume of agent loops can outrun \$/token declines

Conclusion (survey stance)

- Tokens will be a problem** for any product that keeps the LLM in the **online compile loop** without budgets. Treat token burn as OpEx tied to **%gain and p50 wins (C2, P9/P10)**—not as a vanishing cost.
- Inference performance will be a problem** for interactive UX; it is **acceptable** for overnight/CI specialize if SLOs are honest (P18). Do not market batch search as chat.
- Model capabilities will be a problem** in three distinct ways—and they are easy to confuse:
 - **Weak one-shot kernel/IR skill** → needs search + oracles, not a bigger prompt;
 - **Unsafe free rewrite** even when fluent → keep data plane classical (C3/A5);
 - **Fragile tool-calling** on open models → blocks multi-agent SKUs before “coding IQ” does.

4. **Therefore:** under the hybrid prediction, these resource/capability limits **shape the SKU** (freeze, Amdahl trigger, route/distill, FSM bounds) more than they **falsify** agentic compilers. They **do** falsify “LLM-as-opt every build” as a commercial default.
5. **Watch metrics that settle this:** tokens (or \$) per admitted %gain; p50 wall-clock of optimize jobs; tool-call success rate by model tier; win rate after model swap with golden replay (P12/P14).

Example (might be true). A viable commercial control plane spends frontier tokens on **orchestrating a few dozen oracle-gated trials** for Amdahl-hot regions, distills a specialist for the common cases, and ships an ACF/kernel that serves **millions of inferences with zero LLM calls**—so token and capability risk sit in eng CI, not in the customer’s critical path.

Commercial checklist (if you are building this)

- Architecture**
1. Contract = typed tools + structured admit traces (not NL alone).
 2. Memory = scratchpad ≪ dense skills ≪ **VCS artifacts**.
 3. Topology = orchestrator + specialists + shared trace bus; prefer FSM/plan for SLA paths.
 4. Agents run in CI/hot-path → **freeze** before default-on.

- Trust & ops**
5. Layered oracles + named false-negative owners.
 6. CODEOWNERS + signed provenance + sandbox; allowlisted edits + canary rollback.
 7. Joint version pins: agent policy + compiler + HW + model.
 8. Eval ladder: public distributions + private canaries; always show **cost-to-compile**.
 9. A/B / shadow gates before “production default” marketing.
 10. HITL capacity plan (oracle auto-merge narrow; measure review bandwidth).

- Business & compliance**
11. SKU = regressable artifacts (+ CI quota), not chat sessions alone.
 12. Customer owns outputs; telemetry opt-in; pluggable models with golden replay.
 13. Multi-DSL skills; coverage SLA then perf SLA for new silicon.
 14. Tenancy/residency story for ISA RAG and customer graphs.
 15. Support runbooks: replay packs, spend caps, severity for oracle misses.
 16. **Resource envelope (P23):** budget tokens/\$ per %gain; no always-on frontier in default compile; route/distill; measure tool-call success and optimize wall-clock SLOs.

Still blocks / changes packaging: §4.1–4.5, §4.7–4.10; conflicts **C2** (gains pay?), **C3** (API width), **C5** (default-on), **C7** (oracle review), **C9** (coverage vs peak). Resource/capability envelope: **P23**.

5.8 Technical prediction — techniques that accelerate the roadmap

Purpose. §5.5 says *what* should ship by Horizon A/B; §6 says *which checkpoints* can flip the sketch. This subsection answers: **which techniques must be enhanced** so those checkpoints can settle and the roadmap can accelerate — and for each technique, **what is critically missing today**.

Relation to other sections. §4 states gap severity and “done looks like.” §5.7 packages many of the same gaps as commercial problems (P1–P23). Here the cut is **technique-shaped** and **split by locus**: *within the classical compiler / toolchain* versus *outside it* (oracles, data, process, control-plane substrate). Enhancing only opt / Inductor / Triton internals is not enough.

Read rule. Prefer mechanisms over slogans. When a technique’s “missing” column is contested across vendors, keep both sides in §6 rather than averaging.

5.8.1 Within the compiler / toolchain

#	Technique	What exists (illustrative)	Critical missing parts today	Accelerates
T1	Typed agent↔compiler interfaces (tool APIs, structured summaries, enumerated actions)	CompileIQ search surfaces; ACCLAIM/ HintPilot constrained actions; Compiler-R1 tool calls; mlirAgent fingerprints/ MCP; mlir-opt-repl MCP RFC (stateful pass tools); Archer verify/ diffTest; FlashInfer Trace schema; Cake IR (typed schedule + localized findings); Argus tag functions / assertions; TIRx TVM FFI + tile primitives (construct/ inspect/ mutate); TileLang LSP inlay hints (inferred layouts); Gluon / TLX explicit layouts & warp-group ops; Zomboss mapping surface / intent / resolved plan; IntOpt intent sequences; LLM4IR (negative: raw LLVM IR is a poor agent face)	Portable schemas for region / constraints / action / admit across MLIR·Triton·TileLang·CuTe·FlyDSL·Cake·Argus·TIRx; vendor-neutral tool-server conformance (still missing). LLM-oriented IR remains per-stack, not an RFC.	C3, C5, C6; jobs (a)(c); §4.4–4.5

#	Technique	What exists (illustrative)	Critical missing parts today	Accelerates
T2	Admit / fallback machinery (reject illegal agent moves; restore classical path)	Pass-applies-pass patterns (LLM Compiler); template+check admit (AgentCompile); oracle-gated PR review (Archer); TritonRL multi-layer verifiers; FlashInfer-Bench correctness gates before apply(); VibeServe Accuracy Judge; Cake pre-compile safety/conformance gates; Argus layout-algebra + SMT (zero runtime); TIRx well-formed / sync / race / value-sim checks; T-LLM Alive2+CBMC retries; Zomboss Candidate-0 fallback	Shared admit policy as a product surface; deterministic fallback that release eng trusts; open oracles beyond LLVM peephole / single-kernel / NVIDIA-schedule / MI300X-tag / TVM-sim checks	C6 hybrid bet; money-grade shipping; §4.2
T3	Control files, hints, fingerprints + replay	CompileIQ Advanced Control Files; hint/pragma paths; some seed/budget reporting; FlashInfer Trace + apply() as deployable kernel artifacts	Content-addressed cache keys (IR hash, HW, compiler ver, agent policy); golden replay when the agent model upgrades; CI policy for flaky speedups	C2 (p50/p90), C5 (default flag); P4/P14; §4.3
T4	Heuristic hooks & in-tree	Magellan / AlphaEvolve evolve	Settled default between Magellan-class synthesis and MLGO-class advisors on the same apps;	C1 ; Horizon A job (b)

#	Technique	What exists (illustrative)	Critical missing parts today	Accelerates
	advisors (offline job b)	shippable C++; MLGO neural advisors; EmitC-MLGO June 2026 plan-of-record (deploy path advancing)	customer-default EmitC (or public Magellan llvm patches that displace advisors)	
T5	Dialect / ISA feedback sinks (codesign loop into the compiler)	TritonX / KernelEvolve / Ascend diagnosis turn sim+silicon pain into IR/ backend stress; Cake turns recurring failures into verifier rules / IR primitives (corpus-gated); Zomboss compiles TAIDL machine semantics once into ACT backends; sparse RFC narratives	First-class compiler surfaces that aggregate failure modes into dialect or ISA change proposals (still for humans + chip-design tools — not autonomous tape-out)	C9, C10 ; job (d); \$5.5 codesign roadmap

5.8.2 Outside the compiler

#	Technique	What exists (illustrative)	Critical missing parts today	Accelerates
T6	Serving-level oracles & production A/B	Unit/golden/OpInfo; numerical checks; Alive2-class local formal; vendor internal suites; FlashInfer-Bench serving-trace eval + apply() into SGLang/vLLM; VibeServe accuracy/perf judges; GEAK v4 warm-server A/B + output parity + throughput gate; Cake KDA e2e in SGLang	Whole-program / GPU-race / floating-point contracts; multi-month default-path A/B with attribution and public p50/p90; shared open oracles for Triton/Tile beyond FlashInfer operator families	C2 “agents as default”; P5/P20; §4.1–4.2
T7	Open multi-IR corpora (+ negative data)	Meta LLM Compiler (LLVM-heavy); ComPile production LLVM IR; KernelBook (~18k torch↔Triton) → KernelLLM / TritonRL ; DR Triton CSP-DAG synthetic 100k; Compiler-R1; CompilerGym; mostly-closed MLGO corpora	Versioned MLIR / Tile / StableHLO corpora with performance labels and failed compiles / miscompiles / slow-but-correct negatives for critics (Triton positives improved; multi-IR + negatives still thin). ComPile/IRCoder do not close GPU/MLIR gaps.	Selector/Generator quality beyond LLVM-centric models; §4.7
T8	Unified benchmark ladder	Fragmented: CompilerGym, PolyBench/hints, KernelBench(-X); FlashInfer-Bench closes a serving-kernel rung (real traces → leaderboard → deploy); llvm-bench (334 middle-end crash/miscompile bugs); GEAK v4 e2e	Shared ladder IR → single kernel → fused region → full serving graph, reporting correctness × speed × cost-to-compile under fixed HW profiles (FlashInfer-Bench / GEAK e2e / llvm-bench are rungs, not the full ladder)	Honest C2/C9 comparison; kills single-kernel theater; §4.10

#	Technique	What exists (illustrative)	Critical missing parts today	Accelerates
		reports; closed vendor suites		
T9	Provenance, ownership, human-review process	Magellan reviewable C++; Archer oracle review; sparse signing/SBOM discussion; VibeServe git-checkpoint history; SkCC Anti-Skill Injection (untrusted SKILL.md — not kernel CODEOWNERS)	Named CODEOWNERS for agent artifacts; signed admit records (model, tools, oracles, digest); sandbox policy for untrusted proposals; review-capacity metrics	Trusted-base shipping; C7; P6/P16; §4.8–4.9
T10	Agent-workflow compile / freeze / place (control-plane substrate)	FlowCompile, Auto (freeze spans), AgentFlow (ADGs), hetero placement; VibeServe ; DeepSeek Harness (live plugin + session log). Skill plane (2026-08-31): Agent Skills spec; SIGIL ★ SKILL.md → AG-IR → harness; SkCC SkIR + injection; SkVM AOT/JIT; SkillSmith boundary ABI; SKILL.state (_t)	Shared compiler-product agent-graph IR + fail-closed CI (AG-IR/SkIR are skill IRs, not that RFC); placement under spend/latency; kernel VCS still git	Horizon B control-plane compile; P2/P3/P22; §4.6

5.8.3 Checkpoint → technique map

Checkpoint	Needs enhanced techniques	“Settled enough” signal
C1 Magellan vs MLGO	T4 (+ public patches or EmitC customer default)	One path is the default on named shipping apps, or both remain explicitly parallel with release notes
C2 Median / p90 gains	T3 + T6 + T8	Public pinned traces with distributional CI wins and cost-to-compile
C5 Default flag	T1 + T3 (freeze before serve)	Release notes name agent / control-file workflows as supported
C3 / C6 Action width vs replace	T1 + T2	Free rewrite beats advisors on a shared suite or advisory+admit remains the baseline (current lean: hybrid)
C9 Coverage → perf ladder	T5 + bring-up agents + T8	Second-vendor TritonX-class public reproduction
C10 Codesign bound	T5 (proposals only)	Microarch remains human/EDA-owned; agents stress kernels/IR/oracles

5.8.4 Highest-leverage missing parts (near-term)

If an org or research program can fund only a few technique bets to accelerate Horizon A:

1. **Money-grade oracle stack (T2+T6)** — local formal → shape-grid differential tests → statistical serving checks → staged rollout. Without this, agents stay demos.
2. **Replayable artifact contract (T3)** — control files / kernels / heuristics with cache keys and golden replay. Without this, CI rejects the loop.
3. **Portable agent compile interface (T1)** — LLM-oriented IR as summaries · actions · admit records across major AI IRs — **not** a new L-band. Without this, every stack re-implements glue (C3/C4 pressure). Vendor L4 DSLs (TileLang/Gluon/TLX/CuTe/FlyDSL) multiply the glue cost; they do not replace T1.
4. **Open ladder + multi-IR data (T7+T8)** — comparable evaluations and training fuel beyond LLVM-centric or single-kernel suites.

Survey lean. Treat **T1–T5** as compiler/toolchain R&D that must ship as product surfaces (not papers alone), and **T6–T10** as equally first-class — mostly *outside* classical lowering — or the §5.5 roadmap stalls even if model quality improves. Update this subsection when a Tier A / ★ digest closes a “missing” cell or a §6 checkpoint settles.

6. Conflicts (keep unresolved until evidence settles)

This section records **disagreements across papers, vendor blogs, OSS repos, and forums** that matter for predicting the next-generation AI compiler and how agents change that future. We do **not** force a premature resolution; each conflict states both sides, why it matters for the prediction, and what would settle it.

Evidence maps: [../reference/products.md](#) · [../reference/repos.md](#). Technique accelerators for settlement: §5.8.

How to read a conflict row

Field	Meaning
Claim A / Claim B	Competing readings from primary sources
Why it matters	How the winner changes the next-gen compiler sketch
Settlement signal	What evidence would decide it

C1 — Evolve shippable C++ heuristics vs embed neural advisors (Magellan vs MLGO)

Side	Sources	Position
A — Synthesize readable heuristics	Magellan paper/slides; AlphaEvolve lineage; OpenEvolve	Agents rewrite EVOLVE-BLOCK C++ inside LLVM/XLA; ship like human passes; Magellan claims inlining beats decades of manual work; slides hope to leapfrog NN policies
B — Keep/improve neural MLGO	LLVM Discourse EmitC RFC; IR2Vec+MLGO RFC; ongoing MLGO meetings (2026)	Production Chrome/Android/Fuchsia still invest in in-tree NN advisors ; EmitC/TOSA path removes TF build deps so neural policies stay deployable

Why it matters. If A wins, the next-gen control plane is an **offline evolutionary coding agent** whose output is ordinary compiler source. If B wins, the data plane keeps **learned policies** as first-class runtime advisors, and agents mainly help *train/feature* those NNs.

Settlement signal. Public Magellan heuristics land in llvm-project *and* displace MLGO on the same size/perf apps; or EmitC-MLGO becomes the default path for Android/Chrome while Magellan stays Google-internal.

Checkpoint (not settled), June 2026. MLGO sync minutes state a PoR: land Google-internal EmitC for the inliner first, then Android/Fuchsia can start using EmitC; Chrome multi-model support is the next step ([digest](#)). This advances side **B**'s deployability without displacing Magellan or declaring a default.

C2 — Vendor “production agent” wins vs sober benchmark ceilings

Side	Sources	Position
A — Strong commercial wins	NVIDIA CompileIQ blog (Meta up to ~15% on TritonBench/Helion); AMD GEAK v4 (e2e sglang/vLLM A/B) + MLA case study (2.10× / 3.71× TTFT, one MI355 shape); Cake serving-validated KDA 2.05×; Argus 99–104% of MI300X assembly TFLOPS on GEMM/attention/MoE; AlphaEvolve Cloud GA	Agent/autotune control planes already deliver meaningful production speedups
B — Hard ceilings & regressions	CompileIQ docs (often 2–3% on highly optimized kernels); KernelBench / KernelBench-X (many correct kernels slower than eager; refine↑correctness can ↓avg speedup; fusion hard)	Headline speedups are workload-selected; iterative agents can chase correctness at the cost of performance

Why it matters. Prediction of “agents become default compile” needs median/CI wins, not only cherry-picked kernels. Overclaiming delays investment in oracles and traces (§4.2–4.3). [FlashInfer-Bench](#) and [GEAK v4](#) raise serving-trace / warm-server A/B *pressure* (T6/T8) but do **not** settle C2 — still operator- or workload-selected, not multi-month default-path A/B with public p50/p90 + cost-to-compile.

Settlement signal. Reproducible public ACF/kernel agent traces with fixed compiler versions, reporting **distribution** (p50/p90) not only best case; KernelBench-X-style fusion suites remain unsolved or get solved; FlashInfer-class serving traces report distributions under pinned HW/compiler.

C3 — LLMs rewrite IR/code freely vs must stay advisory

Side	Sources	Position
A — Multi-level LLM rewrite works with tests	ACCLAIM (compiler-LLM cooperation); GEAK generate-eval-reflect; KernelAgent; T-LLM (Alive2+CBMC retries)	Guiding agents interleave LLM rewrites with compiler tools; tests/profiles/formal admit candidates; speedups reported
B — Direct IR transform fails	mlirAgent (frontier models below identity on IR transforms); LLM4IR (syntax OK, CFG/exec fail on dumped LLVM IR); HintPilot/AgentCompile design (hints/templates only); Cake (typed schedule beats raw CUDA/PTX); Argus (tag/assert + SMT vs 2-600×-behind free HIP); TIRx (tile primitives + FFI, not dumped TIR text); TileLang LSP (inferred layouts as inspectable facts, not pasted TIR); Zomboss (mapping intents, not native ISA); IntOpt (intent stage, classical realize — still not free opt)	Unconstrained IR/ISA rewrite is unsafe/weak; successful systems constrain the action space

Why it matters. Next-gen architecture either exposes a **wide rewrite API** (with strong oracles) or a **narrow advisory API** (hints, knob ACFs, heuristic blocks). These are different products.

Settlement signal. Shared agent IR contract + oracle suite where free rewrite consistently beats constrained advisors on correctness×perf; or industry standardizes on advisory-only admit gates.

C4 — Kernel DSL future: Triton vs CUDA Tile (and friends)

Side	Sources	Position
A — Triton remains the agent surface	Inductor default path; KernelBench; KernelLLM; GEAK Triton path; awesome-LLM-driven-kernel-generation catalog	Ecosystem, benchmarks, and agents already converge on Triton
B — Tile / CuTe / HIP / FlyDSL / Cake IR / Argus DSL / TIRx / Gluon / TLX fragment the surface	NVIDIA CUDA Tile + CuTe DSL + CompileIQ; TRT-LLM Claude agents for CuTe/TileIR/Triton/CUDA; GEAK multi-language (HIP, FlyDSL, TileLang); Cake IR (explicit schedule, <i>no</i> layout algebra); Argus DSL (layout algebra + SMT tags); TIRx / TileLang (TVM stack); Gluon / TLX (Triton lower/extend — <i>intra</i> -Triton rungs); ThunderKittens (CUDA-embedded)	Vendors/labs push hardware-native tile/schedule IRs; agents must become multi-DSL <i>including</i> Helion vs Gluon vs TLX or lose peak

Why it matters. Training data, tool APIs, and “one agent IR” bets succeed or fail with this choice (§4.4, §4.7).

Settlement signal. One DSL becomes the dominant *agent training* corpus; or a portable tile IR wins; or multi-DSL agents (TRT-LLM skills pattern) become the norm.

C5 — Online compile-time agents vs offline compiler-engineering agents

Side	Sources	Position
A — Online (in the compile/serve loop)	CompileIQ ACFs; HintPilot; AgentCompile; GEAK v4 e2e on sglang/vLLM; Cake kernel evolution; AlphaEvolve Cloud for algo search	Users pay tokens/GPU at optimize time; artifacts are configs/kernels per workload
B — Offline (change the compiler once)	Magellan; MLGO training; Archer PR review; Anthropic Claude C Compiler	Agents change source of the compiler or review PRs; users get classical -O3/opt afterward

Why it matters. These are two different “agent futures.” A hybrid org may need both, but roadmaps and cost models differ.

Settlement signal. Which path shows up as the *default* flag or CI job in PyTorch/LLVM/CUDA release notes over 12–24 months.

C6 — Agents replace compilers vs agents are the control plane

Side	Sources	Position
A — Agents can build/replace large compiler surfaces	Anthropic CCC (~100kLoC Rust compiler); some HN/forum optimism	Agent teams author compilers; classical eng bottleneck shrinks
B — Hybrid control/data plane is the durable pattern	New Compiler Stack survey; mlirAgent limits; ACCLAIM cooperation framing; vendor stacks still ship TRT-LLM/Inductor/XLA	Data plane (lowering, legality, measure) stays classical; agents search/synthesize/advice

Why it matters. Our survey’s executive verdict bets on **B**. A would rewrite goals toward “agent-authored compilers” as the primary object. Timeline of soft merge vs hard replace: §5.1.4 (M1 vs M3).

Settlement signal. A production AI stack whose *default* lowering path is agent-generated without a classical admit/fallback compiler underneath—not a research demo (**M3**). Soft merge of e2e search over classical tools (**M1**) does *not* settle A.

C7 — Generic SCM AI review vs compiler-oracle review

Side	Sources	Position
A — Forge AI is enough	Gerrit ai-code-review / ReviewAI / native AI chat; generic GitHub PR bots	Put LLM on the diff; scale human review
B — Compiler-specialized tools required	Archer (Alive2/LLUBI/opt); llvm-harness / llvm-bench (frontier drop vs SWE-bench; true-fix below 22% after expert review); LLVM Discourse agent-PR experience	Miscompiles need domain oracles; generic review is HITL UX only

Why it matters for *this* survey. Generic Gerrit plugins — and generic coding-agent runtimes such as [DeepSeek Harness](#) or a SKILL .md pack without oracles ([Agent Skills spec](#), skill compilers) — are **weak evidence** for next-gen *compilers*; Archer-class / [llvm-harness](#) tools are strong. Cataloguing forge plugins, star-count harnesses, or SkillsBench pass rates without compiler oracles misaligns with the prediction goal.

Settlement signal. A Gerrit/GitHub bot that blocks merge on failed Alive2/KernelBench-class checks becomes default in llvm-project or a major AI compiler.

C8 — “AI compiler” means DL graph compilers vs LLM-for-LLVM

Side	Sources	Position
A — Compilers for AI models	TVM/XLA/Inductor/TRT-LLM/ OpenVINO product docs	Next-gen = better graph → device stacks (Tile, StableHLO, Neuron NKI)
B — AI for compilers	Magellan, LLM Compiler, Compiler-R1, Archer	Next-gen = agents inside LLVM/ XLA/kernel eng

Why it matters. Commercial catalogs over-weight A; research catalogs over-weight B. Prediction must keep **both stacks converging**, with agents as the cross-cut—not pick one catalog.

Settlement signal. (Already partially here.) Products that expose agent APIs *on* DL compilers (CompileIQ, GEAK, Magellan → XLA) are the convergence proofs.

C9 — Coverage-first bring-up agents vs peak-performance kernel agents

Side	Sources	Position
A — Coverage unlocks the device	TritorX (481 ATen ops, OpInfo, MTIA sim+silicon); KForge on Intel Arc; Zomboss 56/56 verified on Gemmini/PLENA vs 26/56 direct neural	New ASICs need <i>any correct</i> backend before peak kernels; agents should maximize operator coverage first — after a compiled machine contract, not by re-prompting the ISA
B — Perf agents are the product	KernelEvolve, GEAK, AutoKernel, KernelBench fast_p	Without speedups vs eager/ libraries, agentic compile does not pay TCO; coverage-only backends still lose to NVIDIA stacks

Why it matters. The agentic-compiler roadmap needs a **ladder** (coverage → perf) vs a single objective. Codesign programs that only optimize HotGEMM will strand models; programs that only chase OpInfo will never win serving.

Settlement signal. Public playbooks that sequence TritorX-class coverage then KernelEvolve-class perf on the same *shipping* ASIC, with serving metrics — or one objective dominates release criteria industry-wide. Zomboss is an academic analog (research accelerators + sim), **not** that second-vendor reproduction.

C10 — Agentic compiler codesign feedback vs autonomous chip design

Side	Sources	Position
A — Agents co-design silicon	Broad “AI for chip design” narratives; optimism from sim bring-up	LLMs will propose ISA/microarch with compilers in the loop end-to-end
B — Agents stress toolchains; humans/EDA own tape-out	TritonX/KernelEvolve actual scope (kernels, dialects, tests, profilers); this survey’s \$5.5 roadmap	Agentic <i>compilers</i> shorten SW TTM and file ISA/IR pain reports; autonomous tape-out is a different field

Why it matters. Keeps survey focused on the **target agentic compiler**. HW is in scope only when it closes the loop through kernels/IR/oracles.

Settlement signal. A production chip whose microarchitecture was primarily agent-proposed *and* validated via agentic compile oracles — not merely agent-written RTL fragments without compiler admit.

Working stance for this survey (until settlement)

1. Prefer **hybrid control/data plane** (C6-B) as the prediction baseline.
2. Treat Magellan-style **heuristic synthesis** and MLGO **neural advisors** as **parallel production bets** (C1 unresolved).
3. Discount single-number vendor speedups without distribution/oracle context (C2).
4. Assume **constrained actions** + **strong oracles** until free rewrite proves itself (C3).
5. Demote generic SCM AI plugins to Tier C evidence (C7).
6. Keep DL-compiler products as **Tier B baselines**, not as the definition of next-gen (C8).
7. Codesign via **coverage** → **perf agent ladder** on sim+silicon (C9); do **not** expand into autonomous EDA (C10-B).
8. Keep ~6–7 **data-plane bands** (L1–L6 + maturing fleet L7); power/cluster as objectives/ placement/oracles; if bands do not consolidate, ship **pluggable interfaces** (§5.1.1–§5.1.2, A6/S6).
9. Target **e2e-optimal-seeking** architecture: joint / bilevel search under product fitness (F); bands are lower/legality surfaces, not independent greeds (§5.1.3, A7/S7).
10. Treat **soft merge (M1)** of e2e search over band tools as Horizon A–B; do **not** bet that agents **replace** the classical data plane (M3) through ~2031 (C6-B, §5.1.4, A8).

Update this section when a conflict gains a decisive public settlement.

7. Prediction claims ↔ evidence

Living map from falsifiable claims to digests. Update when evidence moves status.

Status: **Supported** · **Contested** · **Watch** · **Falsified**

Architecture (agentic compiler)

ID	Claim	Status	Best evidence	Conflicts
A1	Agents own search/orchestration/synthesis; compilers own lowering, legality, measure, fallback	Supported	ACCLAIM, AgentCompile, HintPilot, mlirAgent (negative), LLM4IR (understanding bound), Cake , Argus , TIRx (FFI + classical tirst pipeline), TLX /Gluon (still Triton-lowered), Zomboss	C3, C6
A2	Four agent jobs stick: (a) online, (b) offline heuristics, (c) engineering/review, (d) bring-up/codesign	Supported	(a) CompileIQ/ GEAK v4 /AutoKernel/Cake/ Argus ; (b) Magellan; (c) Archer/ llvm-harness ; (d) TritonX/KernelEvolve/ Zomboss	C5, C9
A3	ACFs, evolved heuristics, verified kernels, optimization memory, bring-up corpora become first-class artifacts	Supported	CompileIQ, Magellan, KernelBlaster, TritonX, FlashInfer Trace/apply()	—
A4	Defaults stay classical until agents win on <i>distributions</i> in CI	Contested	Vendor blogs vs CompileIQ 2–3% docs, KernelBench-X	C2
A5	Unconstrained LLM will not replace opt/Inductor soon	Supported	mlirAgent; LLM4IR ; hybrid Tier A dominance	C3, C6
A6	Data plane keeps ~6–7 abstraction bands (L1–L6 + maturing fleet L7); agents unify contracts/plugins, not one universal cost model / one IR	Supported (lean)	§5.1.1–5.1.2; ACCLAIM multi-level; mlirAgent; T1/§4.4 plugins	C3, C4, C6
A7	Next-gen targets an e2e-optimal-seeking	Supported (lean)	§5.1.3; ACCLAIM guide+test; T6; C2;	C2, C3, C6

ID	Claim	Status	Best evidence	Conflicts
	architecture: joint/ bilevel search under product fitness (F); bands stay as legality/ lower surfaces; local costs are proposal priors only		AutoKernel Amdahl	
A8	Soft merge (M1): e2e controller over band tools by Horizon A–B; hard replace (M3): agents delete classical admit/ lower — not predicted through ~2031	Supported (lean)	§5.1.4; C6-B; §5.5 “will not ship” LLM-as-opt	C6, C3, C4

Process & stack

ID	Claim	Status	Best evidence	Conflicts
P1	Offline heuristic synthesis and MLGO neural advisors remain parallel bets through 2028	Contested	Magellan vs EmitC-MLGO	C1
P2	Multi-DSL / multi-vendor agent skills become normal	Watch	KForge, GEAK v4, TRT-LLM agents, Helion+CompileIQ, Cake IR, Argus DSL, TIRx/ TileLang, TLX/ Gluon, CuTe DSL, FlyDSL	C4
P3	Compiler-oracle review beats generic forge AI for opt PRs	Supported (direction)	Archer; llvm-harness / llvm-bench; Tier C demoted	C7
S1	Stack reshape is control-plane agentic over classical data plane	Supported	SURVEY §5.6 · A1	C6
S4	Custom ASIC TTM increasingly gated by agentic bring-up	Supported (industrial)	TritorX, KernelEvolve	C9
S5	Profilers/compiler internals become agent APIs	Watch	KernelEvolve MPP, Ascend hierarchical diagnosis, Cake localized diagnostics	C3
S6	Multiple data-plane abstractions remain (~L1–L7); cluster/power attach as place/objectives/oracles; missing consolidation → pluggable interfaces	Supported (lean)	§5.1.1–5.1.2 · A6	C3, C4, C6
S7	Stack reshape centers on e2e fitness (F) + joint controller; freeze only under (F)-admit; siloed per-	Supported (lean)	§5.1.3 · A7	C2, C6

ID	Claim	Status	Best evidence	Conflicts
	band greed is not next-gen			
S8	Soft merge M1 (e2e controller over tools) by Horizon A-B; hard replace M3 not claimed through ~2031	Supported (lean)	§5.1.4 · A8	C6, C3, C4

Codesign (still agentic-compiler-centric)

ID	Claim	Status	Best evidence	Conflicts
H1	Pre-silicon sim + agents provide compiler/ISA feedback before tape-out	Supported (early)	TritorX QEMU future devices	C9, C10
H2	Coverage-first agents then perf agents is the bring-up ladder	Supported	TritorX → KernelEvolve	C9
H3	Agents will not autonomously tape out chips by ~2031; they stress compilers/ISAs	Supported (prediction)	Scope of TritorX/ KernelEvolve (kernels/ toolchains)	C10

Settlement watch

Signal	Moves	Digests
Public Magellan llvm + OpenEvolve recipes	C1	magellan, openevolve
EmitC-MLGO default on Android/ Chrome (PoR advancing: internal inliner → Android/Fuchsia → Chrome multi-model)	C1	mlgo-emitc-rfc, mlgo-emitc-sync-2026-06
p50/p90 public ACF/kernel traces (+ serving-trace distributions)	C2	compileiq-*, kernelbench-x, flashinfer-bench, geek-v4-github
Second non-Meta <i>shipping</i> ASIC reproduces TritorX-class coverage	C9	tritorx (Zomboss = academic analog only)
Agent IR contract where free rewrite beats advisors	C3	acclaim vs hintpilot/ agentcompile, llm4ir, cake, argus
MOCHA / Compiler 2.0 OSS + verified rewrite evals	A1, H1, S4	compiler-2.0-mocha-aarno, compiler-2.0-cgo2026

8. Systems gallery

Snapshot of representative systems. Numbers are **as reported by authors**; cross-benchmark comparison is not apples-to-apples. Prefer mechanisms over headline speedups when choosing architecture.

System	Layer	Agent shape	Correctness gate	Headline	Venue / source
Meta LLM Compiler	LLVM IR / asm	Foundation model	Compiler applies passes	~77% of autotune size potential	CC 2025
Compiler-R1	LLVM pass order	Single agent + tools + RL	opt + IR instr count	~8.5% IR size vs -Oz avg	NeurIPS 2025
LLM-VeriOpt	LLVM IR peephole	Trained small model	Alive2 equivalence	~90% verifiably correct	CGO 2026
Magellan	Heuristics (C++)	Coding agent + AlphaEvolve	Compile + macro-bench	Beats decades of inlining heuristics	C4ML@CGO'26 / LLVM DevMtg
AwareCompiler	Pass sequences	Multi-turn agent-env	Compile/run rewards	Knowledge bridges features↔passes	arXiv 2025
AutoPass	Pass / flags	Multi-agent, inference-only	Compile + profile evidence	Practical budgets, no offline RL	arXiv 2026
HintPilot	Source pragmas	Iterative RAG refine	Compiler validates hints + tests	Up to 6.88× geo-mean vs -Ofast	ACL Findings 2026
ACCLAIM	C / asm multi-level	Guide + level + test agents	LLM tests + compile	Up to 1.25× over compilers alone	arXiv 2026
Reasoning Compiler	TVM schedules	LLM proposals + MCTS	TVM compile/measure	Sample-efficient vs MetaSchedule	NeurIPS 2025
AgentCompile	CUDA inference	Advisor in bounded search	Template CUDA + checks + bench	~4–5.7× vs eager (small LLMs)	arXiv 2026
GEAK	Triton / AMD GPU	Gen/eval/reflect/optimize	Unit tests + timing	Up to 63% exec acc; ~2.59×	AMD / arXiv 2025
KernelLLM	PyTorch → Triton	Specialized 8B model	KernelBench-Triton	Strong Pass@k vs larger general LLMs	Meta HF 2025
mlirAgent	MLIR / LLVM	Tool-using agents (MCP)	Pass tracking + benches	LLMs weak at direct IR rewrite	UCB BAR
LLM4IR / LaMIR	LLVM IR understanding	Probe (no product loop)	CFG / decompile / execute tasks	Syntax OK; CFG/ exec fail	ICML 2025
IntOpt	LLVM IR intent stage	Formulate → refine → realize	Author verified correctness	90.5% correct; 2.66× vs LLM IR gens (author; C2)	arXiv 2026
IRIS-14B	GIMPLE → LLVM IR	Fine-tuned translator	Syntax + I/O tests	Glue layer; opt still classical	arXiv 2026

System	Layer	Agent shape	Correctness gate	Headline	Venue / source
Generative Compilation	Rust frontend	In-decoding feedback	Sealor + rustc	Compiler active during generation	arXiv 2026
Compiler.next	FMware / intent	SE 3.0 vision	Multi-objective quality gates	Compile prompts/agents/params	arXiv 2025
NVIDIA CompileIQ	NVCC/PTXAS knobs	Evolutionary search	Measure on real kernels	$\leq 15\%$ on hot Triton/CUTLASS kernels	CUDA 13.3 blogs
MLGO	LLVM heuristics	RL policies in-tree	Compiler semantics	Prod inlining/regalloc	Google / LLVM
KernelBench	Eval harness	N/A (benchmark)	Correct + fast_p	Frontier <20% one-shot typical	ICML 2025
TritonX	PyTorch ATen / Triton-MTIA	FSM agent bring-up	OpInfo + model ops	481 ops; ~84% pass; sim+silicon	MLSys 2026
KernelEvolve	Triton (+TLX) multi-HW	Graph search + HW RAG	TritonBench + federated profilers	Hetero NVIDIA/AMD/MTIA prod	Meta 2025/26
KForge	Multi-DSL / multi-vendor	Gen $\leftrightarrow$ perf-analysis agents	Compile + correct + profile	+2.12% vs TRT-LLM; 5.13 $\times$ on Arc L2	arXiv 2026
AutoKernel	Triton/CUDA on PyTorch models	Keep/revert agent loop	5-stage harness	Beats eager & torch.compile on hot ops	arXiv 2026
Ascend hierarchical diagnosis	Triton-NPU	Escalating compiler-grounded agents	Profile $\rightarrow$ IR $\rightarrow$ compiler source	4.35 $\times$ geo-mean on 37 ops	arXiv 2026
Helion	PyTorch $\rightarrow$ Triton DSL	Autotune (not LLM)	Config search	Geomean > compile/Triton on reported suites	PyTorch 2025
TIRx	TVM TIR next $\rightarrow$ CUDA C++/PTX	Design: FFI + search rungs (not a public agent loop)	Pre-benchmark wellformed/sync/race/sim	$\sim 0.95\text{--}1.00\times$ library on B200 GEMM/FA4 (expert; C2)	TVM blog 2026
TileLang	TVM tile DSL $\rightarrow$ CUDA/HIP/Metal	Autotune + 2026 LSP (not a public agent loop)	Compile + layout inference / LSP diagnostics	GEMM $\sim$ lib; MHA 1.36 $\times$ vs FA3 (author; C2)	arXiv 2025
Gluon		Hand kernels / experimental	Triton compile		Triton docs 2025/26

System	Layer	Agent shape	Correctness gate	Headline	Venue / source
	Triton ttg IR (layouts, warp specialize)			Blackwell FA-class control (docs; C2)	
TLX	Triton + MIMW extensions	Production kernels; KernelEvolve searches it	ATen-competitive numerical	CUDA-competitive GEMM/attn (author; prod deploy)	arXiv 2026
FlyDSL	Python + Fly MLIR → ROC DL	GEAK/ KernelEvolve sink	ROCm compile	Layout-algebra AMD face (C4)	ROCm blog 2026
CuTe DSL	CUTLASS 4 Python layout algebra	Substrate under ★ CuTeGen	CUTLASS compile + numerical	Near C++ CUTLASS on B200 GEMM/ FMHA (expert; C2)	NVIDIA blog 2025
Event Tensor / ETC	Megakernel IR (events-as-tensors) <i>above</i> L4	Compiler scheduling, not an LLM loop	Static/dynamic task admit	vs vLLM/ SGLang; up to 1.40× TP overlap; ~3.5× warmup (author; C2)	MLSys 2026
ThunderKittens 2.0	CUDA-embedded tiles	“You or your agent” examples, not a published loop	CUDA compile + bench	B200 GEMM ~cuBLAS (expert; C2)	Hazy blog 2026
CuTeGen	CuTe / CUDA	Generate–test–refine + delayed profiling	Compile + numerical + timed	1.71× avg vs PyTorch on KB L1+L2 (author)	arXiv 2026
Cake	Cake IR → CUDA/PTX	Agent + evolving verifier/cost harness	Pre-compile gates + numerical + serving	1.144× vs FlashML clean-start; 2.05× KDA e2e	arXiv 2026
Argus	Argus tile DSL → AMD ISA	ICRL planner + knowledge base	Tag/assert + layout SMT (compile-time)	99–104% of MI300X assembly TFLOPS (author)	arXiv 2026
Zomboss	Generated ASIC backend	Mapping-intent agent	Compiler legalize + verified fallback	56/56 correct; 3.34× / 1.10× vs default	arXiv 2026
GEAK v4	Instinct + sglang/vLLM			Amdahl e2e loop; MLA case	AMD 2026

System	Layer	Agent shape	Correctness gate	Headline	Venue / source
		JS workflow + multi-agent kernel	Warm-server A/B + output parity	2.10× / 3.71× TTFT	
T-LLM Compiler	C loops / LLVM	LLM rewrite + retry	Compiler + Alive2 + CBMC	83.3% author “accuracy” on PolyBench/C	arXiv 2026
llvm-harness	LLVM middle-end	Autofix / autoreview agents	Reproducers + expert review	True-fix below 22% after review; 334-bug bench	arXiv 2026
SIGIL	Skill compile (AG-IR)	SKILL.md → typed harness	Skill-faithfulness (AMC), not Alive2	Mean AMC 66%→88.6%; 2.4–6× fewer tokens (author)	arXiv 2026
SKILL.state	Agent runtime state	Explicit (t); drop reasoning after validate	State-update check ≠ serving (F)	Bounded prompt vs growing transcript (author)	arXiv 2026

Online vs offline agents

Mode	When it runs	Typical artifact	Examples
Online	At compile / specialize time for a program or model	Pass list, hints, kernel choice, ACF knobs	HintPilot, AgentCompile, CompileIQ, Compiler-R1 inference
Offline	Compiler engineering / model training	C++ heuristics, foundation weights, datasets	Magellan, Meta LLM Compiler training, MLGO training

Related engineering experiments (adjacent)

Work	Why it matters to this survey
Anthropic Claude C Compiler	Agents as <i>compiler writers</i> ; harness + tests dominate
LLVM agent PR review Discourse thread	Compiler-specific tools >> generic SWE agents for opt review

9. How to update this survey

Prediction-first loop aimed at the **future agentic compiler** (SW stack + HW codesign feedback). Digests are evidence.

Decision tree

```

New source
├─ Reshapes agentic compile / heuristics / kernels / oracles / ASIC bring-up?
│   YES → Tier A (or B if substrate only: Helion, StableHLO, llvm-project,
│       skill spec / skill compile / \(\Sigma_t\) – P1/P2/T10)
│   NO  → skip or Tier C one-liner
├─ Pure EDA/RTL LLM with no kernel/IR/oracle loop?
│   YES → out of scope (unless it feeds compiler codesign claims H*)
├─ Conflicts with §6 / §7 claims?
│   YES → update §6 Conflicts first (never average)
├─ Closes a §5.8 T* “missing” cell (oracles, schemas, ladder, corpora, ...)?
│   YES → thin-update §5.8 exists/missing + digest ★ if prediction-critical
├─ Moves SURVEY §5 architecture/roadmap/stack/commercial?
│   YES → §7 Claims + narrative; else digest+INDEX+STATUS only

```

Search scope (do not repeat the TIRx miss)

The LLM-oriented IR pass (2026-08-25) searched the **paper cluster titled LLM + compiler IR**. **TIRx** (TVM blog 2026-06-22) was missed until asked: it does not use that phrase. The 2026-08-26 re-research walked **vendor kernel IRs / docs launches / FFI + tile primitives + LSP** and stacks whose last digest was years old. That found TileLang, Gluon, TLX, FlyDSL, CuTe DSL, Event Tensor, ThunderKittens 2.0 — languages SURVEY already *named* in C4/GEAK/KernelEvolve with **zero digests**.

Do	Do not
Vendor blogs + GitHub (Triton/TVM/ROCm/CUTLASS/Hazy)	Only arXiv title search “LLM + IR”
Digest a kernel language if agents treat it as a sink	Digest only the agent paper on top (CuTeGen without CuTe DSL)
Walk stale stacks (TVM froze at TensorIR 2022)	Assume 2022 substrate is current
Disambiguate TritonX ≠ TIRx ≠ TLX	Collapse vendor names

Still undigested (watchlist, not this wave): JAX Pallas / Mosaic GPU; **HipKittens** (AMD TK); Modular Mojo GPU kernels. Same rule if they become named agent sinks.

Search scope (do not repeat the SKILL.state miss)

The post-TIRx walks (2026-08-26) searched **vendor kernel IRs**. **SKILL.state** (arXiv:2608.26263, posted ~2026-08-25) was missed until asked: the title has no compiler/IR, and in this repo “agent skills” had meant CompileIQ / TensorRT-LLM **compile packs**. The public stack also has (a) the **Agent Skills** SKILL.md spec and (b) ****skill compilation / (_t)**** papers. Those are **P1 / P2 / T10 substrate**, not a fifth compiler job and not a new L-band.

Do	Do not
Walk SKILL.md, agentskills.io, “skill compilation,” AG-IR / SkIR / SkVM, “agent execution state”	Only treat “skills” as CompileIQ / TRT-LLM packs
Digest the spec if SURVEY names SKILL.md as the pack format	Digest only the vendor pack that follows it
Keep three senses distinct (compile pack ≠ spec ≠ skill compile / (t))	Collapse homonyms; cite “SkillSmith” without an arXiv number
Mount a compiler control plane or an oracle loop → Tier B substrate	Generic SWE agents, customer-service skills, CTF / exploit write-ups as compile evidence

Still undigested (skill watchlist, Tier B only): EvoSkill (skill *discovery*, QA benches); Formal Skill / FairyClaw (runtime-native JSON+hooks, not SKILL.md compile); SSL (skill-text → structure, 2604.24026); Agent skill security (2607.13987); SkillSmith *co-evolve* (2606.01314 — **not** the compile paper 2605.15215). Digest when a compiler/oracle mount appears.

Add-source order

1. Tier A/B/C
2. Digest from reference/publications/_TEMPLATE.md (create if missing; fill **Org** + **Publisher**)
3. INDEX row with Org/Publisher columns (★ only for prediction-critical)
4. §6 Conflicts if disagreeing
5. §7 Claims if prediction moves
6. Thin touch to SURVEY §5 (incl. §5.8 T-table when technique evidence lands)
7. reference/repos.md / reference/products.md / §8 Systems if mechanism new
8. STATUS changelog
9. python3 scripts/validate_survey.py → git push origin main

Depth

Stub → digest → deep (PDF) before changing SURVEY §5.5 horizons.

Reference guide

Evidence store for the living survey. Digests, product signals, and repo/forge maps live here; the **prediction narrative** stays in `../docs/SURVEY.md` (§0–§9).

Start here

Category	Path	What it is
Publications	<code>publications/ · INDEX.md</code>	One digest per searched source (papers, blogs, talks, code pages). ★ = prediction-critical.
Products	<code>products.md</code>	Commercial SKUs / offerings as prediction signals (Tier A/B/C) — not a full catalog.
Repos	<code>repos.md</code>	GitHub / Gerrit / forge artifacts tiered for the agentic-compiler prediction — not an exhaustive forge list.

How to use

1. Read `../docs/SURVEY.md` §0 → §5 for the prediction (roadmap §5.5, stack §5.6, commercial §5.7).
2. Open ★ digests from `publications/INDEX.md` when you need mechanism detail.
3. Use `products.md` / `repos.md` for Tier A/B/C shipping surfaces (CompileIQ, GEAK, ACCLAIM, Archer, ...).
4. When sources disagree → `../docs/SURVEY.md` §6. Claim IDs → §7. Systems snapshot → §8.

Add a source

1. Create a digest from `publications/_TEMPLATE.md` (**Org** + **Publisher** required).
2. Add an INDEX row (★ only if it moves the prediction).
3. Update `products.md` / `repos.md` only if the mechanism is a new shipping surface.
4. Update SURVEY §6 / §7 if a durable claim or disagreement moved; thin-touch §5 if the prediction moves.
5. Follow `../docs/SURVEY.md` §9 (How to update this survey). For “skills”: walk the `SKILL.md` spec, skill compilers, and (`_t`) runtimes (hard rule 15) — not only vendor compile packs.
6. `python3 scripts/validate_survey.py`

Do **not** paste long digests into `docs/SURVEY.md`. Keep the narrative thin; park evidence here. The survey points into this tree.

Related narrative

- `../docs/SURVEY.md` — single reading path: vocabulary (§0), Q1–Q4, prediction (§5), conflicts (§6), claims (§7), systems (§8), update loop (§9)

Products (prediction signals)

Purpose: Company offerings that inform **what next-gen compile will ship** and **how agents enter production** — not a full SKU catalog.

Companion: [../docs/SURVEY.md §5](#) · [../docs/SURVEY.md §6](#) · [repos.md](#) · [guide](#)

Role tags

Tag	Meaning
DataPlane-Compile / Serve	Graph → device or inference engine (baseline agents must plug into)
ControlPlane-Autotune	Searches knobs/schedules (classic or evolutionary)
ControlPlane-Agent	Explicit LLM/agent in the loop
Kernel-Platform	Tile/DSL/IR agents target
Cloud-Agent	Sold as cloud agent service
Research-release	Important, not a full SKU

Tier A — shapes the agent future

Company	Offering	Roles	Prediction signal	Conflicts
Google Cloud / DeepMind	AlphaEvolve on Cloud (GA)	Cloud-Agent, ControlPlane-Agent	Evolutionary coding agent as a sold product; Magellan lineage	C5, C1
Google	Magellan (prod inlining narrative) + MLGO	ControlPlane-Agent / Autotune	Offline heuristic synthesis <i>and</i> neural advisors in production	C1
NVIDIA	CompileIQ (+ agent-skills)	ControlPlane-Autotune / ControlPlane-Agent	Workload-specialized compiler controls; versioned ACFs; AGENTS.md skill pack drives search+Welch validate	C2 (blog vs docs)
NVIDIA	CUDA Tile / Tile IR	Kernel-Platform	Next agent IR vs Triton	C4
AMD	GEAK v4 (e2e_workflow + kernel_workflow)	ControlPlane-Agent, Kernel-Platform	Amdahl triage on warm sglang/vLLM; warm-server A/B + output parity; multi-DSL kernels	C2, C4, T6
Meta	LLM Compiler / KernelLLM	Research-release	Open foundation / specialist models	—
NVIDIA	TensorRT-LLM + Claude agents/skills PR	DataPlane-Serve + ControlPlane-Agent	Agents wired into flagship serve compiler (multi-DSL)	C4
Meta	TritonX + KernelEvolve (MTIA + hetero GPUs)	ControlPlane-Agent + Kernel-Platform	Agentic ASIC bring-up + production ranking kernels	C9, C2
Meta / LF	Helion	Kernel-Platform	Higher-level agent/autotune surface over Triton	C4
FlashInfer / NVIDIA·UW·CMU	FlashInfer-Bench (+ Trace / apply())	ControlPlane-Agent + DataPlane-Serve	Serving-trace kernel ladder; dynamic substitution into SGLang/vLLM	C2, T6/T8
NVIDIA · CMU	Cake (Cake IR; FlashInfer PRs)	ControlPlane-Agent, Kernel-Platform	Typed schedule IR + evolving verifier; serving-validated KDA / TinyGEMM	C3, C4, T1/T5

Company	Offering	Roles	Prediction signal	Conflicts
CausalFlow et al.	Argus (paper; MI300X)	ControlPlane-Agent, Kernel-Platform	Data-flow invariants + SMT; 99–104% of assembly TFLOPS (author)	C2, C3, C4, T1/T2

Tier B — data-plane baselines (agents must integrate)

Brief on purpose: these are **defaults**, not proof that agents win.

Company	Offering	Roles	Why keep
NVIDIA	TensorRT-LLM / CUDA libs	DataPlane-Serve/Compile	Peak production path
FlashInfer	flashinfer kernel library	DataPlane-Serve	Engine-agnostic attention/GEMM/MoE kernels; Bench deploy target
Meta	<code>torch.compile</code> / Inductor	DataPlane-Compile	De-facto app compile entry
Google / OpenXLA	XLA + StableHLO	DataPlane-Compile	Portable HLO; Magellan XLA experiments
Modular	MAX + Mojo	DataPlane-Serve/Compile	MLIR-rooted alternative stack
Intel	OpenVINO	DataPlane-Compile	Edge/CPU deploy toolkit
AWS	Neuron (+ NKI)	DataPlane-Compile, Kernel-Platform	Cloud-custom silicon; NKI as agent surface
Qualcomm	AI Hub / QNN	Edge compile	On-device compile-as-a-service
Qualcomm	Hexagon-MLIR	DataPlane-Compile, Kernel-Platform	Open Triton/PyTorch → Hexagon NPU MLIR stack

Demoted / footnote (misaligned with prediction goal)

Item	Why demoted
Olive / ORT / HF Optimum glue	Cross-runtime packaging; little agent-compile signal
OctoML (historical TVM SaaS)	Cite only as lineage for commercial autotune appetite that later products (CompileIQ, AlphaEvolve Cloud) still sell — not an active next-gen agent compiler
Generic “AI code review” SKUs	Conflict C7 — HITL UX, not compiler oracles
Anthropic CCC	Process signal (agents-as-engineers), not a sold compiler SKU
DeepSeek Harness (dsh)	Open coding-agent runtime (plugin kernel + session log). C7 unless compiler oracles mount as seams; T10 color only — digest
Agent Skills spec (SKILL.md)	Packaging format (progressive disclosure). C7 unless compiler oracles mount; CompileIQ already follows it — digest
SIGIL / SkCC / SkVM / SkillSmith / SKILL.state	Skill <i>compilation</i> or (<code>_t</code>) runtimes. T10/P2 substrate, not a sold compiler SKU — see publications INDEX

Architecture placement

Control plane (emerging SKUs)
AlphaEvolve Cloud · GEAk v4 · CompileIQ · Magellan/MLGO · TRT-LLM agent skills · Cake IR
Data plane (mature defaults)
TRT-LLM · Inductor · XLA · MAX · OpenVINO · Neuron
Kernel / tile platforms
Triton · Helion · Gluon · TLX · CUDA Tile · CuTe DSL · TileLang · TIRx · FlyDSL · ThunderKittens · NKI · HIP

Commercial reality: revenue still sits on the data plane. Explicit agent control-plane SKUs are newer and often opt-in for hot kernels — consistent with \$5 prediction (defaults classical first).

Mapping to survey questions

Question	Commercial signal
Q1 Trends — hybrid control plane	CompileIQ, GEAK v4, Cake, AlphaEvolve Cloud, Magellan/MLGO
§1b Traditional still wins defaults	TRT-LLM, Inductor, XLA, OpenVINO
Q2 How agents help	Evolve code, kernel loops, knob search, specialist models
Q3 Reshape process	ACF-in-VCS; heuristic synthesis; agent skills in TRT-LLM
§5 Future	Tier A rows above + conflicts C1–C5

Commercial digests for Tier A/B prediction-relevant items live under [publications/](#) (CompileIQ, GEAK, AlphaEvolve, Magellan/MLGO, TRT-LLM agents). Skip Olive/OctoML churn.

Update when a vendor ships a **named agent-compile default**, not when another runtime EP appears.

Repos & forge evidence

Purpose: Map GitHub / Gerrit / googlesource artifacts to the survey **prediction** (next-gen compiler + how agents change the future). This is **not** an exhaustive forge catalog.

Companion: [../docs/SURVEY.md §5](#) · [../docs/SURVEY.md §6](#) · [products.md](#) · [guide](#)

Evidence tiers

Tier	Meaning	Use in prediction
A — Reshapes compile	Agents change heuristics, kernels, knobs, or compiler review with domain oracles	Primary evidence for §5
B — Substrate	Data-plane hosts agents must address	Required context; not “agent future” alone
C — Delivery / HITL only	Generic forge AI or tangential hosting	Demoted; cite only for conflict C7

Tier A — agents reshape compilation

Repository	Forge	Future role	Gaps / conflicts
cuhk-s3/Archer + paper	GitHub	Compiler-oracle PR review (Alive2/LLUBI/opt)	§4.2, §4.8; conflict C7
algorithmicsuperintelligence/openevolve	GitHub	OSS AlphaEvolve-style loop (Magellan OSS path)	§4.1, C1
cornell-zhang/heurigym	GitHub	Agentic heuristic bench incl. compiler tasks	§4.10
amazon-science/acclaim + paper	GitHub	Multi-level compiler↔LLM cooperation (online job a)	Q2/Q3; §5.4
ucb-bar/mlirAgent	GitHub	MCP + fingerprints; negative free-IR-rewrite result	§4.5; C3
hjiang13/LLM4IR	GitHub	IR <i>understanding</i> probe (CFG/exec); not a product compiler	C3 ; T1
Mind4Compiler/Compiler-R1	GitHub	Tool-using RL pass agent	Q2
ZJU-PL/hintpilot	GitHub	Constrained hint/pragma synthesis	C3 advisory path
ScalingIntelligence/KernelBench	GitHub	Kernel LLM benchmark	Trend D; C2
BonnieW05/KernelBenchX (if public) / paper	GitHub	Correctness*perf ceilings	C2
meta-pytorch/KernelAgent	GitHub	PyTorch→verified Triton agents	Trend D
AMD-AGI/GEAK / GEAK-agent	GitHub	v4: e2e sglang/vLLM Amdahl + A/B; multi-DSL kernel_workflow	Trend D; C2/C4 ; T6
dtxzyw/llvm-harness + paper	GitHub	LLVM tools + llvm-bench + autofix/autoreview (Archer lineage)	Job (c); C7
NVIDIA/CompileIQ	GitHub	Evolutionary compiler Advanced Controls → ACF	§4.3; C2/C5
anthropics/claude-c-compiler	GitHub	Agents-as-compiler-engineers	Trend F; C6
flagos-ai/awesome-LLM-driven-kernel-generation	GitHub	Living kernel-agent bibliography	Trend D watchlist

Repository	Forge	Future role	Gaps / conflicts
RightNow-AI/autokernel	GitHub	Amdahl-driven model-level kernel agent loop	Job (a); C2
TritonX / KernelEvolve (papers; code mostly internal)	Meta	ASIC bring-up + hetero perf agents (MTIA/ NVIDIA/AMD)	Job (d); C9
flashinfer-ai/flashinfer-bench	GitHub	Serving-trace kernel ladder + apply() into SGLang/vLLM	T6/T8; C2
uw-syfi/vibe-serve	GitHub	Agentic end-to-end serving-stack synthesis	T6/T10 ; research
pytorch/helion	GitHub	High-level tile DSL → Triton (agent/CompileIQ surface)	Tier B substrate; C4

Magellan itself remains mostly internal (paper + [LLVM Dev Meeting slides](#)); track OSS via OpenEvolve + HeuriGym until a public Magellan tree appears (**C1**). Slides next-steps: OpenEvolve OSS path; XLA green-field / auto-sharding — folded into [.. /docs/SURVEY.md](#) §5.4.

Tier B — substrate (data plane hosts)

Repository	Why it matters for prediction
llvm/llvm-project	Host for MLGO, Magellan heuristics, Archer reviews
google/ml-compiler-opt	Neural advisor training stack (parallel bet to Magellan)
facebookresearch/CompilerGym	RL pass-order gym — substrate for Selector agents, not agent future alone
apache/tvm (tvm.tirx)	TIRx kernel IR + FFI mutation surface (T1/C4 ; digest)
tile-ai/tilelang + tilelang-lsp	TVM tile DSL + layout-inference LSP (T1/C4 ; digest)
mlc-ai/tirx-kernels	Community TIRx kernel library / B200 benches
triton-lang/triton	Default GPU DSL; Gluon is the ttg lower face
facebookexperimental/triton	TLX MIMW extensions (KernelEvolve sink)
ROCm/FlyDSL	AMD Python+MLIR layout DSL (C4 ; digest)
NVIDIA/cutlass	CuTe DSL Python (CUTLASS 4); substrate under CuTeGen
HazyResearch/ThunderKittens	CUDA-embedded tiles (C4 ; digest)
openxla/xla / StableHLO	Portable AI compiler IR; Magellan XLA experiments
PyTorch (torch.compile / Inductor)	Default DL compile path agents must plug into
NVIDIA/TensorRT-LLM	Production serve stack; PR #12831 adds Claude kernel/compile agents (C4)
flashinfer-ai/flashinfer	Kernel library / generator backing many serve engines; Bench <code>apply()</code> target
Hugging Face GPUMODE/KernelBook	Open PyTorch↔Triton training pairs (T7)
deepseek-ai/deepseek-harness	Agent <i>runtime</i> (Cordis plugins + session log). T10 substrate; C7 — not a compiler harness (digest)
agentskills/agentskills	SKILL.md spec (digest). Packaging, not a compiler
sigilagent/sigil	SIGIL: SKILL.md → AG-IR → harness (T10 ★ ; digest)
Nexa-Language/Skill-Compiler	SkCC / SkIR emitters (digest)
AetherHeart-AI/Aeloon	SkillSmith boundary compile (digest)

Tier C — demoted (delivery channel only)

These show demand to put LLMs **inside** code review UX. They are **not** next-gen compiler semantics unless wired to Alive2/opt/KernelBench-class oracles (conflict **C7**).

Project	Link	Note
Gerrit ai-code-review	googlesource	Generic diff LLM
ReviewAI Gerrit plugin	amarula/reviewai-gerrit-plugin	Sidebar chat
GerritForge AI provider	GerritForge/ai-review-agent-provider	Interface layer only

Open opportunity (still Tier A if built): Gerrit/GitHub plugin that **calls compiler oracles** before commenting — Archer pattern on Gerrit.

Implications for the predicted future

1. **SCM is part of the control plane** only when oracles are present (Archer), not when a chatbot comments on a diff (Tier C).
2. **Offline heuristic evolution** (OpenEvolve/Magellan) and **online knob/kernel agents** (CompileIQ/GEAK) are both Tier A — different jobs (C5).
3. **Negative results are Tier A too** — mlirAgent's below-identity IR rewrite bounds the architecture (§5.1).
4. Prefer Tier A/B when enriching digests; do not grow Tier C catalogs.

Digests: [publications/INDEX.md](#).

Publications index

Publications index

One digest per searched source. Files live beside this index. **Org** = company/university/lab;
Publisher = venue, blog host, forge, or preprint host.

Year	Kind	Group	Org	Publisher	Digest	Source
2026	paper	Surveys & vision	ICT, CAS · UCAS · University of Leeds	arXiv	The New Compiler Stack: A Survey on the Synergy of LLMs and Compilers	source
2025	paper	Surveys & vision	Queen's University	arXiv	Compiler.next: A Search-Based Compiler to Power the AI-Native Future of Software Engineering	source
2026	paper	Surveys & vision	Qualcomm Technologies International	arXiv	Reading AI Model Compilation in MLIR Through the Lens of Formal Theories	source
2026	talk	Surveys & vision	MIT CSAIL	ACM/IEEE Ken Kennedy Award · HPCA/CGO/PPoPP/CC 2026	Compiler 2.0: Building the Next Generation Compilers with ML (Ken Kennedy Award plenary, HPCA/CGO/PPoPP/CC 2026) ★ vision	source
2022	talk	Surveys & vision	MIT CSAIL	CGO 2022 · YouTube	CGO 2022 Keynote: Compiler 2.0	source
2020	talk	Surveys & vision	MIT CSAIL	ACM CC 2020	Compiler 2.0: Using ML to Modernize Compiler Technology (CC 2020)	source
2025–28	company	Surveys & vision	Aarno Labs · MIT · UIUC (DARPA MOCHA)	Aarno Labs / DARPA	DARPA MOCHA / Aarno Labs: Compiler 2.0 project ★ funded path	source
2018	paper	Classic DL compilers	UW · AWS · Berkeley et al.	USENIX OSDI 2018	TVM: An Automated End-to-End Optimizing Compiler for Deep Learning	source
2020	paper	Classic DL compilers	UW · AWS · OctoML et al.	USENIX OSDI 2020	Ansor: Generating High-Performance Tensor Programs for Deep Learning	source
2021	company		Apache TVM			source

Year	Kind	Group	Org	Publisher	Digest	Source
		Classic DL compilers		Apache TVM blog	TVM blog: Introducing Auto-scheduler (Anso)	
2020	paper	Classic DL compilers	Peking University et al.	ACM ASPLOS 2020	FlexTensor: An Automatic Schedule Exploration and Optimization Framework	source
2022	paper	Classic DL compilers	UW · AWS · OctoML et al.	ACM ASPLOS 2023	TensorIR: An Abstraction for Automatic Tensorized Program Optimization	source
2026	company	Classic DL compilers	Apache TVM · MLC (CMU course)	Apache TVM blog	TIRx: An Open Compiler Stack for Evolving Frontier ML Kernels	source
2025/26	paper	Classic DL compilers	Peking University · Microsoft Research · Imperial College London	arXiv	TileLang: A Composable Tiled Programming Model for AI Systems	source
2021	paper	Classic DL compilers	Google · multi-institution	arXiv (MLIR)	MLIR: A Compiler Infrastructure for the End of Moore's Law	source
2025	company	Classic DL compilers	OpenXLA / Google	OpenXLA	OpenXLA StableHLO roadmap	source
2021	paper	MLGO & RL gyms	Google · CMU	arXiv	MLGO: a Machine Learning Guided Compiler Optimizations Framework	source
2022	company	MLGO & RL gyms	Google Research	Google Research blog	Google Research blog: MLGO	source
2022	company	MLGO & RL gyms	InfoQ (covering Google MLGO)	InfoQ	InfoQ: MLGO Framework Brings Machine Learning in Compiler Optimizations	source
2022+	code	MLGO & RL gyms	LLVM / Google MLGO	LLVM docs	LLVM docs: Machine Learning	source

Year	Kind	Group	Org	Publisher	Digest	Source
					Guided Optimization (MLGO)	
2022	code	MLGO & RL gyms	Meta FAIR	GitHub	CompilerGym (Meta/FAIR)	source
2022	forum	MLGO & RL gyms	LLVM community (GSoC)	LLVM Discourse	LLVM Discourse: ML-guided ordering of compiler optimization passes (GSoC)	source
2023	paper	Foundation LLMs for compilers	Meta	arXiv	Large Language Models for Compiler Optimization	source
2024	paper	Foundation LLMs for compilers	Meta	arXiv	Meta Large Language Model Compiler: Foundation Models of Compiler Optimization	source
2024	company	Foundation LLMs for compilers	Meta AI	Meta AI Research	Meta AI publication page: LLM Compiler	source
2024	company	Foundation LLMs for compilers	Meta (Chris Cummins)	LinkedIn	Chris Cummins: LLM Compiler foundation models post	source
2024	paper	Foundation LLMs for compilers	Meta	arXiv	Compiler generated feedback for Large Language Models	source
2023	forum	Foundation LLMs for compilers	Hacker News community	Hacker News	HN: Large Language Models for Compiler Optimization	source
2024	forum	Foundation LLMs for compilers	Hacker News community	Hacker News	HN: Meta LLM Compiler (long thread)	source
2024	forum	Foundation LLMs for compilers	Hacker News community	Hacker News	HN: Meta Large Language Model Compiler	source
2023	paper	Foundation LLMs for compilers	UC Davis · TUM · LLNL · UIUC ·	arXiv · Hugging Face	ComPILE: A Large IR Dataset from Production Sources	source

Year	Kind	Group	Org	Publisher	Digest	Source
			ANL · Google et al.			
2024	paper	Foundation LLMs for compilers	TU Darmstadt · University of Würzburg	ACL 2024 · arXiv	IRCoder: Intermediate Representations Make Language Models Robust Multilingual Code Generators	source
2025	paper	Foundation LLMs for compilers	Kent State · HUST · PNNL · Chongqing University	ICML 2025 · PMLR · arXiv	Can Large Language Models Understand Intermediate Representations in Compilers? ★ IR understanding bound	source
2026	paper	Foundation LLMs for compilers	Barcelona Supercomputing Center · UPC	arXiv	LLM Translation of Compiler Intermediate Representation	source
2025	paper	Agentic & RL compilers	ISCAS · UCAS	arXiv	Compiler-R1: Towards Agentic Compiler Auto-tuning with Reinforcement Learning	source
2025	code	Agentic & RL compilers	ISCAS / Mind4Compiler	GitHub	Mind4Compiler/Compiler-R1 (code)	source
2026	paper	Agentic & RL compilers	multi-institution	CGO 2026	LLM-VeriOpt: Verification-Guided RL for LLM-Based Compiler Optimization	source
2026	paper	Agentic & RL compilers	Google DeepMind / Google	arXiv	Magellan: Autonomous Discovery of Novel Compiler Optimization Heuristics with AlphaEvolve ★ Tier A offline job	source
2025	talk	Agentic & RL compilers	Google DeepMind / Google	LLVM Developers' Meeting 2025	LLVM Developers' Meeting slides:	source

Year	Kind	Group	Org	Publisher	Digest	Source
					Magellan ★ future signals (\$5.4)	
2025	paper	Agentic & RL compilers	Google DeepMind	arXiv	AlphaEvolve: A coding agent for scientific and algorithmic discovery	source
2026	company	Agentic & RL compilers	Google DeepMind	DeepMind blog	DeepMind blog: AlphaEvolve impact	source
2025	paper	Agentic & RL compilers	ISCAS · UCAS · NTU Singapore	arXiv	AwareCompiler: Agentic Context-Aware Compiler Optimization	source
2026	paper	Agentic & RL compilers	Shaanxi Normal University · Northwest University · University of Leeds	arXiv	AutoPass: Evidence-Guided LLM Agents for Compiler Performance Tuning	source
2026	paper	Agentic & RL compilers	SKLP, ICT, CAS · UCAS · Jiangnan University · CUHK	arXiv	Beyond Pass-by-Pass Optimization: Intent-Driven IR Optimization with Large Language Models	source
2026	paper	Agentic & RL compilers	Zhejiang University · Purdue	arXiv	HintPilot: LLM-based Compiler Hint Synthesis for Code Optimization	source
2026	paper	Agentic & RL compilers	AWS AI · Georgia Tech	arXiv	Agentic Code Optimization via Compiler-LLM Cooperation (ACCLAIM) ★ Tier A Q2/Q3	source
2026	paper	Agentic & RL compilers	Huawei Technologies, Heterogeneous Compiler Lab	arXiv	T-LLM Compiler: Trusted LLM-based Code Optimization and Verification Framework	source
2026	code	Agentic & RL compilers	Amazon Science / AWS AI	GitHub	amazon-science/acclaim (GitHub) ★ Tier A	source
2026	paper		ETH Zurich · INSAIT/Sofia	arXiv	Generative Compilation: On-	source

Year	Kind	Group	Org	Publisher	Digest	Source
		Agentic & RL compilers	University · UC Berkeley		the-Fly Compiler Feedback as AI Generates Code	
2025	paper	GPU kernels & inference compilers	Stanford · Princeton	arXiv	KernelBench: Can LLMs Write Efficient GPU Kernels?	source
2025	company	GPU kernels & inference compilers	Stanford (Scaling Intelligence)	Stanford blog	Stanford blog: KernelBench	source
2025	code	GPU kernels & inference compilers	Stanford (Scaling Intelligence)	GitHub	ScalingIntelligence/KernelBench	source
2026	paper	GPU kernels & inference compilers	Tsinghua University	arXiv	KernelBench-X: A Comprehensive Benchmark for Evaluating LLM-Generated GPU Kernels	source
2025	paper	GPU kernels & inference compilers	AMD	arXiv	GEAK: Introducing Triton Kernel AI Agent & Evaluation Benchmarks	source
2025	company	GPU kernels & inference compilers	AMD	AMD ROCm blog	AMD ROCm blog: Triton kernel AI / GEAK	source
2025	code	GPU kernels & inference compilers	AMD AGI	GitHub	AMD-AGI/GEAK agent	source
2025	company	GPU kernels & inference compilers	Meta	Hugging Face	Meta KernelLLM (Hugging Face)	source
2025	company	GPU kernels & inference compilers	GPU MODE	Hugging Face	GPUMODE/KernelBook dataset ★ T7 corpus	source
2025	paper	GPU kernels & inference compilers	Amazon · multi-institution	arXiv	TritonRL: Training LLMs to Think and Code Triton Without Cheating	source

Year	Kind	Group	Org	Publisher	Digest	Source
2026	paper	GPU kernels & inference compilers	Texas A&M · Oracle	arXiv	DR Triton: Large-Scale Synthetic Data Driven RL for Triton	source
2026	paper	GPU kernels & inference compilers	UW · CMU · NVIDIA · UC Berkeley	MLSys 2026	FlashInfer-Bench: Virtuous Cycle for AI-driven LLM Systems ★ T6/T8 serving ladder	source
2026	code	GPU kernels & inference compilers	FlashInfer / NVIDIA · UW · CMU	GitHub	flashinfer-ai/flashinfer-bench ★	source
2025	paper	GPU kernels & inference compilers	UC San Diego	arXiv	Reasoning Compiler: LLM-Guided Optimizations for Efficient Model Serving	source
2026	paper	GPU kernels & inference compilers	City University of Hong Kong	arXiv	AgentCompile: An LLM-Guided Compiler for Direct CUDA Inference	source
2025	code	GPU kernels & inference compilers	UC Berkeley (ucb-bar)	GitHub	ucb-bar/mlirAgent	source
2025	company	Company infra	NVIDIA	NVIDIA Developer blog	NVIDIA: Focus on Your Algorithm—CUDA Tile Handles the Hardware	source
2026	company	Company infra	NVIDIA	NVIDIA Developer blog	NVIDIA: Develop High-Performance GPU Kernels in C++ with CUDA Tile	source
2026	company	Company infra	NVIDIA	NVIDIA Developer blog	NVIDIA CUDA 13.3: Tile C++ + CompileIQ	source
2026	company	Company infra	NVIDIA	NVIDIA Developer blog	NVIDIA: Extract More Kernel Performance with CompileIQ	source
2026	code	Company infra	NVIDIA	NVIDIA docs	CompileIQ documentation	source
2025	company		Modular	Modular blog		source

Year	Kind	Group	Org	Publisher	Digest	Source
		Company infra			Modular: What about the MLIR compiler infrastructure?	
2026	company	Company infra	Anthropic	Anthropic Engineering	Anthropic: Building a C compiler with a team of parallel Claudes	source
2026	company	Company infra	Ars Technica (covering Anthropic)	Ars Technica	Ars Technica: Sixteen Claude agents created a C compiler	source
2026	company	Company infra	Modular	Modular blog	Modular/Lattner: The Claude C Compiler — Future of Software	source
2026	forum	Company infra	Hacker News community	Hacker News	HN: We tasked Opus 4.6 agent teams to build a C Compiler	source
2025	forum	Forums & workshops	LLVM community	LLVM Discourse	LLVM Discourse: LLVM ♥ ML Workshop 2025 agenda	source
2026	forum	Forums & workshops	LLVM community	LLVM Discourse	LLVM Discourse: LLVM ♥ ML Workshop 2026 CFP	source
2023	forum	Forums & workshops	LLVM community	LLVM Discourse	LLVM Discourse: ML-Guided Compiler Optimization Workshop 2023	source
2026	forum	Forums & workshops	LLVM community	LLVM Discourse	LLVM Discourse: Automated review with agents (~30 bugs / 207 PRs)	source
2026	forum	Forums & workshops	comp.compilers community	comp.compilers	comp.compilers: Magellan paper notice	source
2025	company	Forums & workshops	IEEE EMBS Pulse	IEEE Pulse	IEEE Pulse: LLMs in Compiler Optimization — Challenges and Future Direction	source

Year	Kind	Group	Org	Publisher	Digest	Source
2026	company	Forums & workshops	Moonlight	Moonlight	Moonlight literature review: Magellan	source
2010s+	code	Correctness lineage	Google	GitHub	Souper: A superoptimizer for LLVM IR	source

2026 | paper | Source control & review agents | CUHK (cuhk-s3) | arXiv | [Archer: Towards Agentic Review for Compiler Optimizations](#) | [source](#) |

2026 | paper | Source control & review agents | SUSTech · ETH Zurich · CUHK | arXiv | [Agentic Harness for Real-World Compilers \(llvm-harness\)](#) ★ job (c) harness | [source](#) |

2026 | code | Source control & review agents | CUHK (cuhk-s3) | GitHub | [cuhk-s3/Archer \(GitHub\)](#) | [source](#) |

2024+ | code | Source control & review agents | Google | Gerrit googlesource | [Gerrit plugin: ai-code-review \(googlesource\)](#) | [source](#) |

2025+ | code | Source control & review agents | Amarula Solutions | GitHub | [amarula/reviewai-gerrit-plugin](#) | [source](#) |

2025+ | code | Source control & review agents | GerritForge | GitHub | [GerritForge/ai-review-agent-provider](#) | [source](#) |

2025 | code | Source control & review agents | Algorithmic SuperIntelligence | GitHub | [OpenEvolve \(AlphaEvolve-style OSS\)](#) | [source](#) |

2025 | code | Source control & review agents | Cornell University | GitHub | [HeuriGym \(cornell-zhang/heurigym\)](#) | [source](#) |

2025 | code | Source control & review agents | Meta (PyTorch) | GitHub | [meta-pytorch/KernelAgent](#) | [source](#) |

2026 | code | Source control & review agents | NVIDIA | GitHub | [NVIDIA/CompileIQ \(GitHub\)](#) | [source](#) |

2026 | code | Source control & review agents | Anthropic | GitHub | [anthropics/claudes-c-compiler](#) | [source](#) |

2021+ | code | Source control & review agents | Google | GitHub | [google/ml-compiler-opt](#) | [source](#) |

2026 | code | Source control & review agents | Zhejiang University (ZJU-PL) | GitHub | [ZJU-PL/hintpilot](#) | [source](#) |

ongoing | code | Source control & review agents | LLVM Foundation / community | GitHub | [llvm/llvm-project \(host repository\)](#) | [source](#) |

2026 | paper | Surveys & vision | BAAI · PKU · HKUST(GZ) et al. | arXiv | [Towards Automated Kernel Generation in the Era of LLMs](#) | [source](#) |

2025+ | code | GPU kernels & inference compilers | FlagOpen / flagos-ai | GitHub | [flagos-ai/awesome-LLM-driven-kernel-generation](#) | [source](#) |

2026 | company | Commercial products & proposals | AMD | AMD ROCm blog | [AMD ROCm: GEAK v3 kernel optimization agent](#) | [source](#) |

2026 | code | Commercial products & proposals | AMD | GitHub | [GEAK v4: end-to-end GPU performance optimization system](#) ★ e2e serving A/B | [source](#) |

2026 | company | Commercial products & proposals | AMD | AMD ROCm blog | [GEAK Agent-Driven Optimization of the DeepSeekV4 MLA Kernel](#) | [source](#) |

2025 | forum | Forums & workshops | Google / LLVM community | LLVM Discourse | [LLVM](#)

[Discourse RFC: EmitC support for MLGO](#) | [source](#) |
 2026 | forum | Forums & workshops | Google / LLVM community | LLVM Discourse | [LLVM MLGO sync minutes: EmitC path PoR \(June 8, 2026\)](#) ★ C1 checkpoint | [source](#) |
 2026 | forum | Forums & workshops | LLVM / MLIR community | LLVM Discourse | [LLVM Discourse RFC: mlir-opt-repl + MCP server](#) | [source](#) |
 2026 | code | Commercial products & proposals | NVIDIA | GitHub (TensorRT-LLM) | [TensorRT-LLM PR: Claude agents/skills for kernels and compile](#) | [source](#) |
 2026 | company | Commercial products & proposals | NVIDIA | NVIDIA docs | [CompileIQ docs: expected gains \(2 to 3 percent highly optimized\)](#) | [source](#) |
 2026 | company | Commercial products & proposals | NVIDIA | NVIDIA CompileIQ docs | [CompileIQ agent-skills \(AGENTS.md optimization campaigns\)](#) ★ control-plane UX | [source](#) |
 2025/26 | paper | HW codesign & accelerator bring-up | Meta | arXiv | [Agentic Operator Generation for ML ASICs \(TritonX\)](#) ★ bring-up / codesign | [source](#) |
 2025/26 | paper | HW codesign & accelerator bring-up | Meta | ISCA 2026 · arXiv | [KernelEvolve: Scaling Agentic Kernel Coding for Heterogeneous AI Accelerators at Meta](#) ★ multi-HW prod | [source](#) |
 2026 | company | HW codesign & accelerator bring-up | Meta | Engineering at Meta | [Meta Engineering blog: KernelEvolve / Ranking Engineer Agent](#) | [source](#) |
 2026 | paper | HW codesign & accelerator bring-up | Huawei Technologies | arXiv | [Compiler-Grounded Hierarchical Diagnosis for Triton on NPUs](#) ★ Ascend | [source](#) |
 2026 | paper | HW codesign & accelerator bring-up | University of Michigan · The University of Texas at Austin | arXiv | [Rethinking Agentic Kernel Generation for Emerging Accelerators \(Zomboss\)](#) ★ compile-once mapping contract | [source](#) |
 2026 | paper | HW codesign & accelerator bring-up | Gimlet Labs | MLArchSys @ ISCA 2026 · arXiv | [KForge: Cross-Platform Kernel Generation for AI Accelerators](#) ★ multi-DSL | [source](#) |
 2026 | paper | GPU kernels & inference compilers | RightNow AI | arXiv | [AutoKernel: Autonomous GPU Kernel Optimization](#) ★ Amdahl agent loop | [source](#) |
 2026 | code | GPU kernels & inference compilers | RightNow AI | GitHub | [RightNow-AI/autokernel \(GitHub\)](#) ★ | [source](#) |
 2026 | paper | GPU kernels & inference compilers | University of Michigan | arXiv | [Kernel Forge: Agent Harness for CUDA Kernel Opt](#) | [source](#) |
 2026 | paper | GPU kernels & inference compilers | University of Toronto · Standard Kernel | arXiv | [CuTeGen: Agentic GPU Kernels using CuTe](#) ★ CuTe lane / C4 | [source](#) |
 2026 | paper | GPU kernels & inference compilers | NVIDIA · UC Berkeley | arXiv | [KernelBlaster: Memory-Augmented In-Context RL for CUDA](#) | [source](#) |
 2026 | paper | GPU kernels & inference compilers | NVIDIA · Carnegie Mellon University | arXiv | [CAKE: Compiler-Agent Co-Design for Frontier Kernel Evolution](#) ★ typed schedule IR | [source](#) |
 2026 | paper | GPU kernels & inference compilers | CausalFlow · HKUST · Tsinghua University · Stanford · UCAS · UC Riverside | arXiv | [ARGUS: Agentic GPU Optimization Guided by Data-Flow Invariants](#) ★ SMT/layout admit | [source](#) |
 2025/26 | company | GPU kernels & inference compilers | Triton project (OpenAI / community) | Triton docs | [Gluon: Triton's lower-level GPU programming model](#) | [source](#) |
 2026 | paper | GPU kernels & inference compilers | UC San Diego · Meta | arXiv | [TLX: Hardware-Native, Evolvable MIMW GPU Compiler for Large-scale Production Environments](#) | [source](#) |
 2026 | company | GPU kernels & inference compilers | AMD | AMD ROCm blog | [FlyDSL: Expert GPU Kernel Development with the Ease of MLIR Python Native DSL on AMD GPUs](#) | [source](#) |
 2025 | company | GPU kernels & inference compilers | NVIDIA | NVIDIA Developer blog | [Achieve CUTLASS C++ Performance with Python APIs Using CuTe DSL](#) | [source](#) |

2026 | paper | GPU kernels & inference compilers | Carnegie Mellon University · NVIDIA et al. | MLSys 2026 · arXiv | [Event Tensor: A Unified Abstraction for Compiling Dynamic Megakernel](#) | [source](#) |

2026 | company | GPU kernels & inference compilers | Stanford Hazy Research | Hazy Research blog | [ThunderKittens 2.0: Even Faster Kernels for Your GPUs](#) | [source](#) |

2025 | company | Classic DL compilers | Meta (PyTorch) | PyTorch blog | [Helion: High-Level DSL for Portable ML Kernels](#) | [source](#) |

2026 | paper | Classic DL compilers | Qualcomm | arXiv | [Hexagon-MLIR: AI Compilation Stack for Hexagon NPUs](#) | [source](#) |

2025+ | code | Classic DL compilers | Meta (PyTorch) | GitHub | [pytorch/helion \(GitHub\)](#) | [source](#) |

2026 | paper | Agent control-plane substrate | RightNow AI | arXiv | [Auto: The AGI Compiler](#) ★ [freeze/amortize](#) | [source](#) |

2026 | paper | Agent control-plane substrate | UMass Amherst · MIT · MIT-IBM Watson AI Lab | arXiv | [FlowCompile: An Optimizing Compiler for Structured LLM Workflows](#) ★ [workflow compile](#) | [source](#) |

2026 | paper | Agent control-plane substrate | Huazhong University of Science and Technology · Macquarie University | arXiv | [AgentFlow: Building Agent Dependency Graphs for Static Analysis of Agent Programs](#) | [source](#) |

2025 | paper | Agent control-plane substrate | Stanford · Gimlet Labs · Intel | arXiv | [Efficient and Scalable Agentic AI with Heterogeneous Systems](#) | [source](#) |

2026 | paper | Agent control-plane substrate | University of Washington (SyFI) | arXiv | [VibeServe: Can AI Agents Build Bespoke LLM Serving Systems?](#) | [source](#) |

2026 | code | Agent control-plane substrate | DeepSeek AI | GitHub | [DeepSeek Harness \(dsh\)](#) [plugin runtime / session log](#) | [source](#) |

2026 | paper | Agent control-plane substrate | Google LLC · Purdue University | arXiv | [SKILL.state: Scalable Long-Horizon Agent Skills](#) | [source](#) |

2025–26 | company | Agent control-plane substrate | Anthropic (origin) · agentskills.io ecosystem | [agentskills.io / GitHub](#) · Anthropic Engineering | [Agent Skills \(SKILL.md\)](#) [open specification](#) | [source](#) |

2026 | paper | Agent control-plane substrate | University of Michigan | arXiv | [SIGIL: Compiling Agent Skills into Typed Harnesses](#) ★ [skill compilation](#) | [source](#) |

2026 | paper | Agent control-plane substrate | Sun Yat-sen University | arXiv · ACM CAIS 2026 (AgentSkills'26) | [SkCC: Portable and Secure Skill Compilation for Cross-Framework LLM Agents](#) | [source](#) |

2026 | paper | Agent control-plane substrate | Shanghai Jiao Tong University | arXiv | [SkVM: Compiling Skills for Efficient Execution Everywhere](#) | [source](#) |

2026 | paper | Agent control-plane substrate | AetherHeart Tech · Renmin University of China · UC San Diego | arXiv | [SkillSmith: Compiling Agent Skills into Boundary-Guided Runtime Interfaces](#) | [source](#) |

Total: 142 digests

Kinds: paper · company · forum · talk · code

★ = high-signal for next-gen **prediction** (prefer these when updating [docs/SURVEY.md](#) §5).

See also: [reference guide](#) · [repos.md](#) / [products.md](#) (Tier A/B/C) · [docs/SURVEY.md](#) §6 (conflicts) · [docs/SURVEY.md](#) §5.6.

Export notes

- Full digests remain in `reference/publications/*.md` (not inlined; each has Org + Publisher).
- Reference: `reference/README.md` → `publications / products / repos`.
- Commercialization critical problems: `docs/SURVEY.md` §5.7 (P1–P23).
- Technical prediction (techniques to accelerate roadmap): `docs/SURVEY.md` §5.8.
- Rebuild: `python3 publish/build_pdf.py`.